

15. Joint distributions: The autoregressive time-series model y_1, y_2, \dots with mean level 0, autocorrelation 0.8, residual standard deviation 1, and normal errors can be written as $(y_t | y_{t-1}, y_{t-2}, \dots) \sim N(0.8y_{t-1}, 1)$ for all t .

- (a) Prove that the distribution of y_t , given the observations at all other integer time points t , depends only on y_{t-1} and y_{t+1} .
 (b) What is the distribution of y_t given y_{t-1} and y_{t+1} ?

$$(a) \quad y_t \sim \frac{1}{\sqrt{2\pi}} \exp\left\{-\frac{1}{2}(y_t - 0.8y_{t-1})^2\right\}$$

$$P(A|B) = \frac{P(A, B)}{P(B)}$$

$$P(Z|A, B, \dots) = \frac{P(Z, A, B, \dots)}{P(A, B, \dots)}$$

what is $P(y_{t-1}, y_{t-2}, \dots)$?

I don't know what this is as I am sure that these are not independent y_{t-1}, y_{t-2}, \dots

because

$$f(y_1, \dots, y_t, y_{t+1}, \dots, y_n) = \prod_{t=1}^n N(y_t | 0.8y_{t-1}, 1)$$

$$\Rightarrow f(y_t | y_{t-1}) \propto N(y_t | 0.8y_{t-1}, 1) N(y_{t+1} | 0.8y_t, 1)$$

main idea: y_t is expressed in terms of y_{t+1} and y_{t-1} .

we can continue to say that r.v. y_t is proportional to a distribution that relies on y_{t+1} and y_{t-1} .

$$(b) \Rightarrow f(y_t | y_{t-1}, y_{t+1}) \propto N(y_t | 0.8y_{t-1}, 1) N(y_{t+1} | 0.8y_t, 1)$$

$$\propto \frac{1}{\sqrt{2\pi}} e^{-\frac{1}{2}(y_t - 0.8y_{t-1})^2} \frac{1}{\sqrt{2\pi}} e^{-\frac{1}{2}(y_{t+1} - 0.8y_t)^2}$$

$$= \frac{1}{2\pi} \exp\left\{-\frac{1}{2}\left[y_t^2 - 1.6y_{t-1}y_t + 0.64y_{t-1}^2 + y_{t+1}^2 - 1.6y_{t+1}y_t + 0.64y_t^2\right]\right\}$$

here's a magic trick here! It's literally completing the square.

$$\propto \frac{1}{2\pi} \exp\left\{-\frac{1}{2}\left[1.64y_t^2 - 1.6y_{t-1}y_t - 1.6y_{t+1}y_t\right]\right\}$$

what distn is this?

$$\propto \frac{1}{2\pi} \exp\left\{-\frac{1}{2}\left[1.64y_t^2 - 1.6(y_{t-1} + y_{t+1})y_t\right]\right\}$$

$$\propto \frac{1}{2\pi} \exp\left\{-\frac{1}{2}\left[y_t^2 - \frac{1.6}{1.64}(y_{t-1} + y_{t+1})y_t\right]\right\}$$

$$\left(-\frac{1}{2} \cdot \frac{1.6}{1.64}(y_{t-1} + y_{t+1})\right)^2$$

set to

$$C_1 = \left(\frac{0.8}{1.64}\right)^2 (y_{t-1} + y_{t+1})^2$$

$$\propto \frac{1}{2\pi} \exp\left\{-\frac{1}{2}\left[y_t^2 - 2\sqrt{C_1}y_t + C_1 - C_1\right]\right\}$$

$$\propto \frac{1}{2\pi} \exp\left\{-\frac{1}{2}\left[(y_t - \sqrt{C_1})^2 - C_1\right]\right\}$$

$$\propto \frac{1}{2\pi} \exp\left\{-\frac{1}{2}(y_t - \sqrt{C_1})^2\right\}$$

this is some const

$$= \frac{1}{2\pi} \exp\left\{-\frac{1}{2}(y_t - \sqrt{C_1})^2\right\}$$

$$= N(y_t | \sqrt{C_1}, 1)$$

this distribution is $y_t \sim N\left(\frac{0.8}{1.64}(y_{t-1} + y_{t+1}), 1\right)$

with

$$E[y_t] = \left(\frac{0.8}{1.64}\right)(y_{t-1} + y_{t+1})$$

$$\text{Var}(y_t) = 1$$

Homework 3

Michael Pena

2024-10-27

Book Question 3.11

```
####(a)
# Bioassay
J = 4
n = c(5,5,5,5)
y = c(0, 1, 3, 5)
x = c(-0.86, -0.30, -0.05, 0.73)

#  $Y_i | \theta_i \sim \text{Bin}(n_i, \theta_i)$ ;  $\text{logit}(\theta_i) = \alpha + \beta x_i$ 
#  $p(\alpha, \beta) = 1$ 

Prior = function(alpha, beta){
  dmvnorm(x = c(alpha,beta), mean = c(0,10), sigma = matrix(c(4,10,10,100),2,2))
}

Likelihood = function(Y, N, X, alpha, beta){
  th = inv.logit(alpha + beta*X)
  prod(dbinom(Y, N, th))
}

rProposal = function(n, alpha, beta, s, S){
  rmvnorm(1, mean = c(alpha, beta), sigma = s*S)
}

### initialize
accept = 0
s = 1 # tuning param.
S = matrix(c(1, 0, 0, 1), nrow = 2, ncol = 2)

#  $S = \text{matrix}(c(1, 2.3, 2.3, 18), \text{nrow} = 2, \text{ncol} = 2)$ 
###

temp = y/n # "MLE's of thetas"
temp = c(0.01, 0.2, 0.6, 0.9) # fixing the boundaries
ifit = summary(lm(logit(temp) ~ x)) # est. are reasonable est. for alpha and beta

alpha.post = beta.post = numeric()
alpha.post[1] = -0.3356453
beta.post[1] = 4.2419651

B = 8000
```

```

for(b in 2:(B+1)){

  #Joint MH step: alpha, beta/data
  ab.star = rProposal(1, alpha.post[(b-1)], beta.post[(b-1)], s, S)

  # Ratio of Densities
  r.num = Likelihood(y, n, x, ab.star[1], ab.star[2]) *
    Prior(ab.star[1], ab.star[2])

  r.den = Likelihood(y, n, x, alpha.post[(b - 1)], beta.post[(b - 1)]) *
    Prior(alpha.post[(b - 1)], beta.post[(b - 1)])

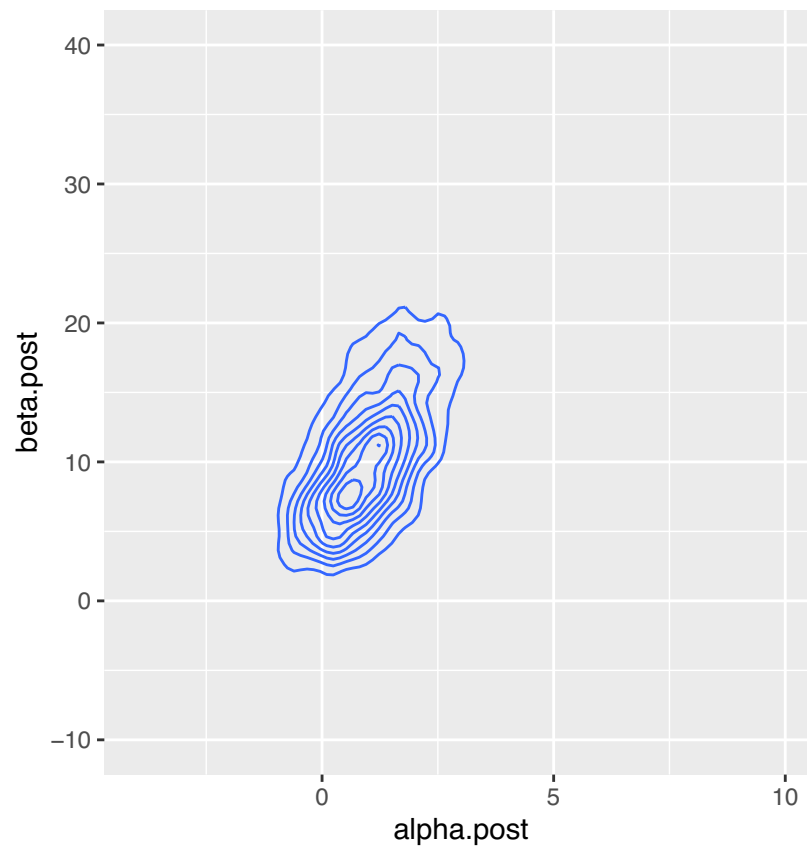
  r = r.num / r.den

  # Accept/Reject
  if(runif(1) < min(1, r)){
    alpha.post[b] = ab.star[1]
    beta.post[b] = ab.star[2]
    accept = accept + 1
  }else{
    alpha.post[b] = alpha.post[(b-1)]
    beta.post[b] = beta.post[(b-1)]
  }
}

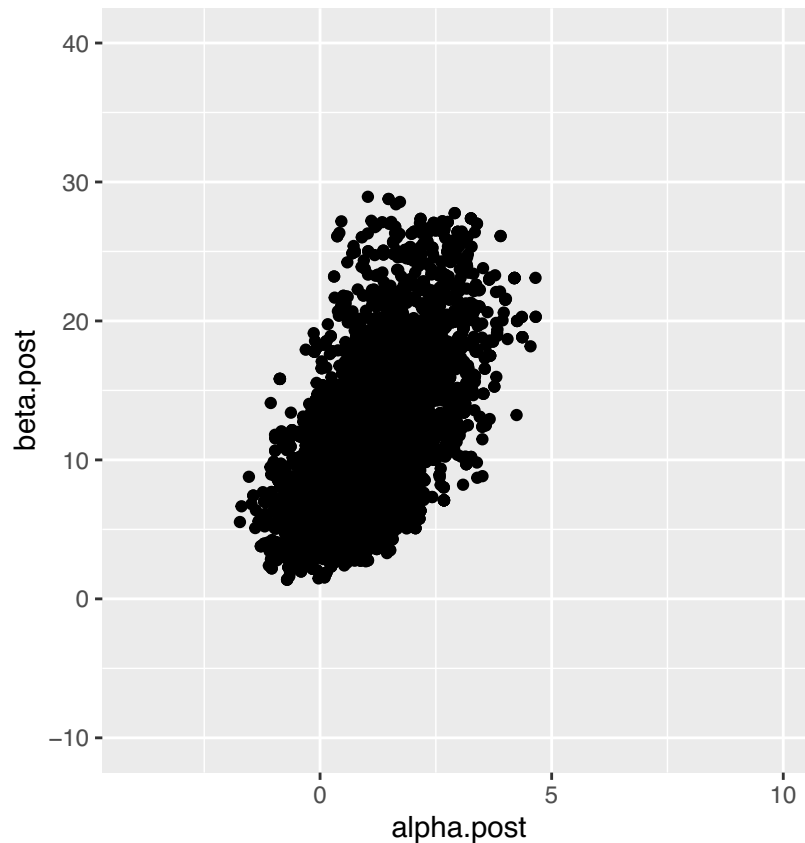
# plot(alpha.post, type = "l")
# plot(beta.post, type = "l")
# burn = 500
# acf(alpha.post[-c(1:burn)])
# acf(beta.post[-c(1:burn)])
#
# var(alpha.post[-c(1:burn)])
# cov(alpha.post[-c(1:burn)], beta.post[-c(1:burn)])
# var(beta.post[-c(1:burn)])

# contour and dots
post <- data.frame(cbind(alpha.post,beta.post))
ggplot(post, aes(x = alpha.post, y =beta.post)) + geom_density2d() +ylim(-10,40) +xlim(-4,10) +coord_fixed()

```



```
ggplot(post, aes(x = alpha.post, y =beta.post)) + geom_point()+ylim(-10,40) +xlim(-4,10)+coord_fixed(ratio=1)
```



! Caption for the image

Our new priors seem to align our scatter and contour plot with that of the previous uniform priors. Perhaps there is more of a variance towards the top of the ellipse compared to when we use the normal to sample our priors. Applying the context, our prior assumptions do not matter much, as we will continue to predict the same amount of deaths depending on log-dosage and number of animals.

Question 1

####(a)

input data (y-vector)

determine recognize that the posterior can be separate as so.

$$P(\mu, \sigma^2 | data) = P(\mu | \sigma^2, data) \cdot P(\sigma^2 | data)$$

Generate σ^2 from $Sc.Inv - \chi^2(n-1, s^2)$

use σ^2 to generate μ from $N(\bar{y}, \frac{\sigma^2}{n})$

do the last two steps B amount of times

####(b)

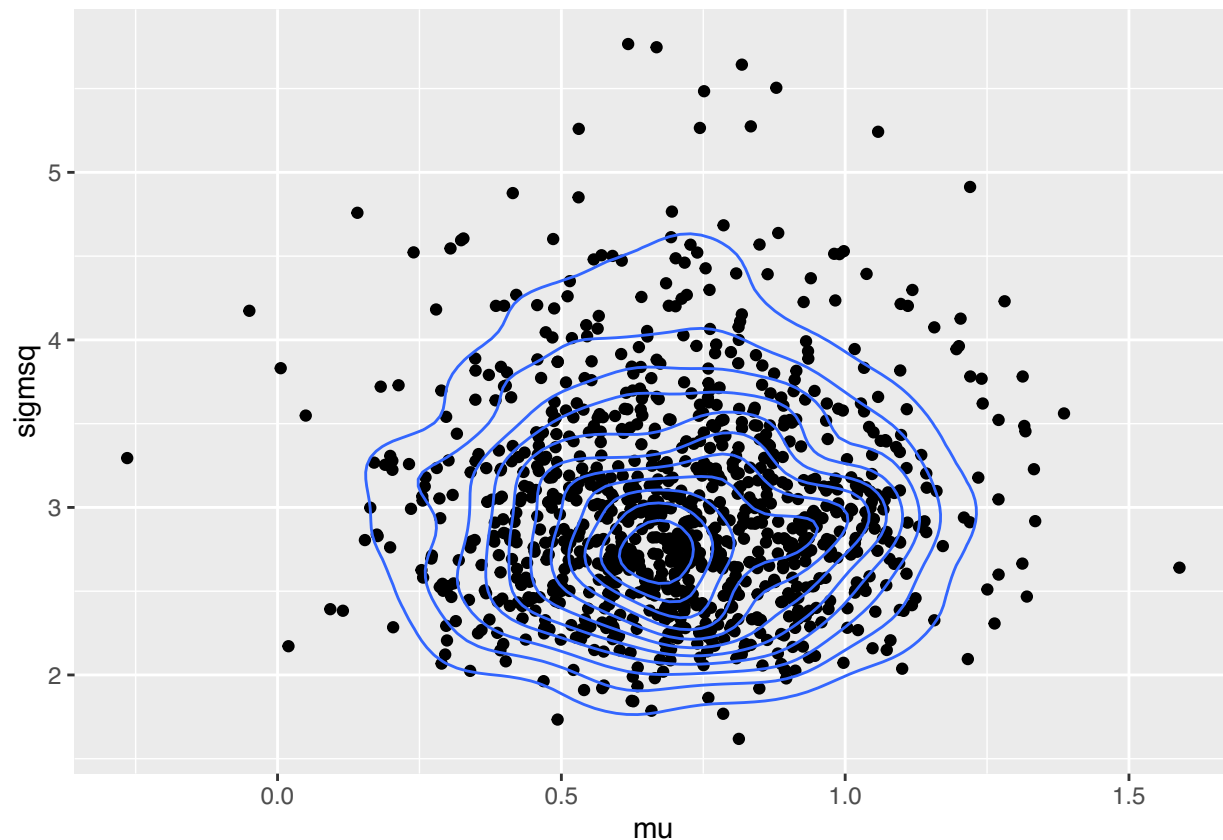
```
set.seed(88899)
# grab some data
n = 50
y = rnorm(n, 1, 2)
ybar = mean(y)
```

```

sy = sd(y)
# grab sigma-squareds from scaled inverse chisquared
B = 1000
sigmsq <- (n-1) * sy^2/rchisq(B,n-1)
# grab mus from normal
mu <- rnorm(B,ybar,sqrt(sigmsq/n))

post <- data.frame(cbind(mu,sigmsq))
ggplot(post, aes(x = mu, y =sigmsq)) + geom_point() + geom_density2d()

```

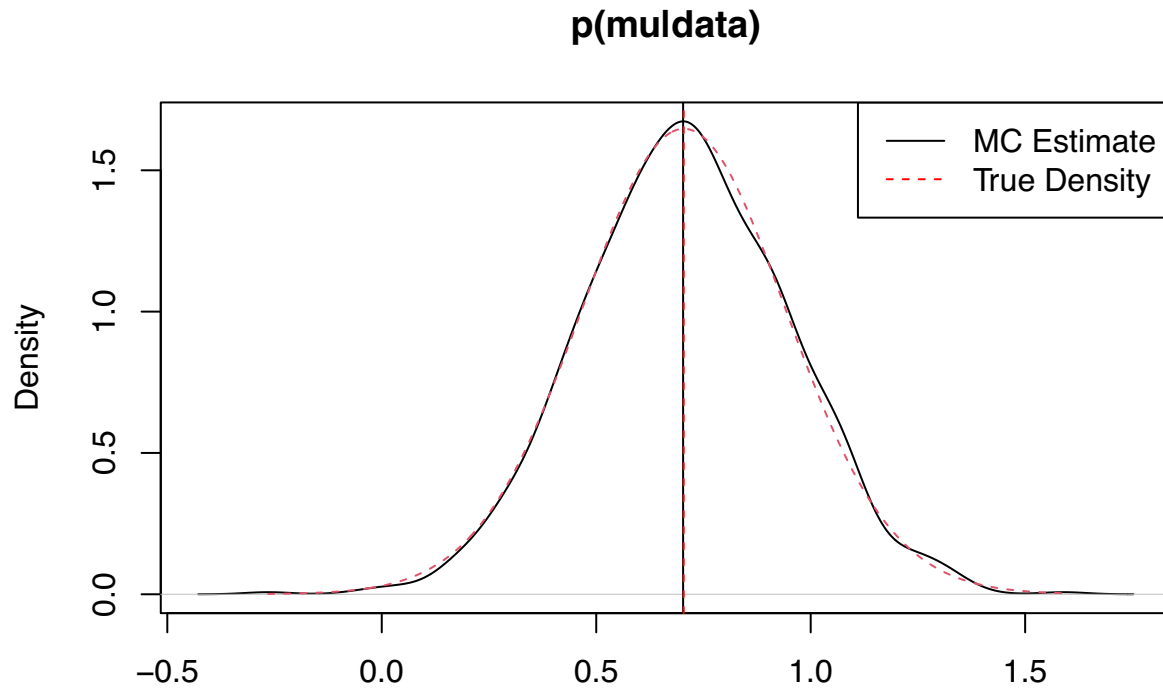


```
####(c)
```

```

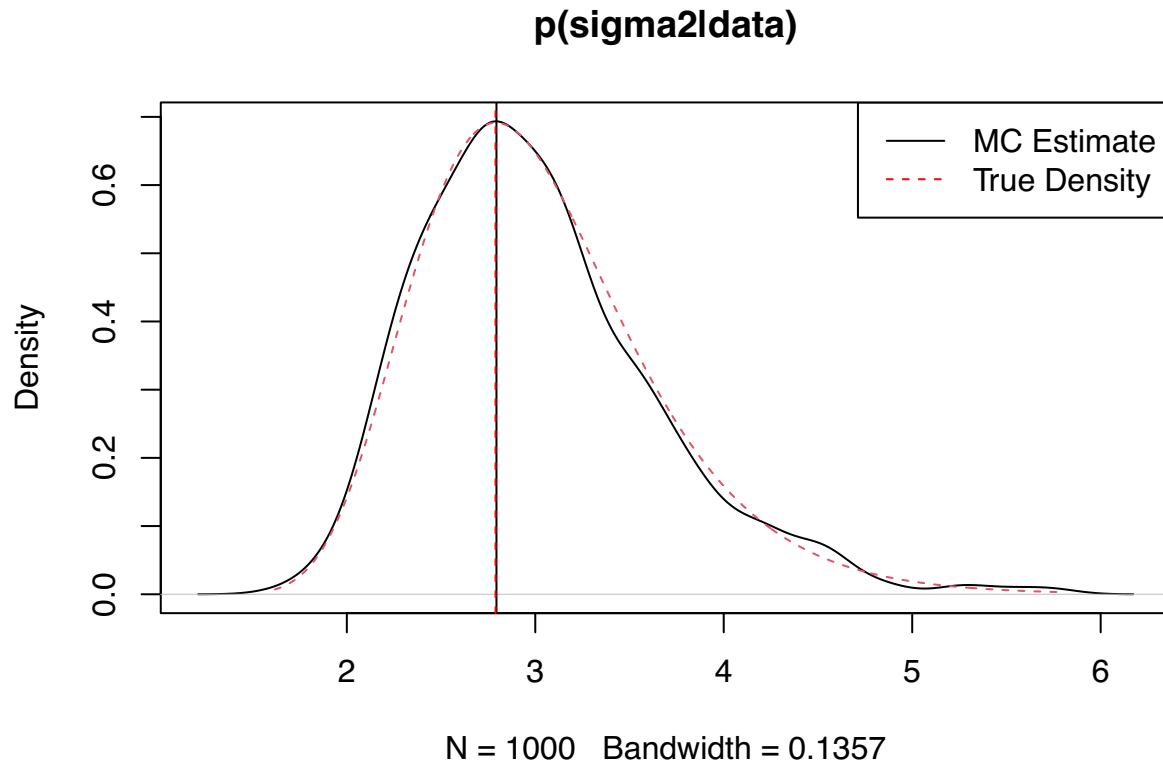
# Plot Marginal Posterior of mu
plot(density(mu), main = "p(mu|data)")
mu.seq = seq(min(mu), max(mu), length.out = B)
lines(mu.seq, dt.scaled(mu.seq, n-1, ybar, sqrt(sy^2/n)), col = 2, lty = 2)
legend("topright", legend = c("MC Estimate", "True Density"), col = c("black", "red"), lty = c(1, 2))
# get the MAP value
density(mu)$x[density(mu)$y == max(density(mu)$y)] -> v1
abline(v=v1)
# get actual MLE
av1 = mu.seq[dt.scaled(mu.seq, n-1, ybar, sqrt(sy^2/n)) == max(dt.scaled(mu.seq, n-1, ybar, sqrt(sy^2/n)))]
abline(v = av1, col = "red", lty = 2)

```



N = 1000 Bandwidth = 0.05438

```
# Plot Marginal Posterior of sigma
plot(density(sigmsq), main = "p(sigma2|data)")
sigmsq.seq = seq(min(sigmsq), max(sigmsq), length.out = B)
lines(sigmsq.seq, dinvchisq(sigmsq.seq, n-1, sy^2), col = 2, lty = 2)
legend("topright", legend = c("MC Estimate", "True Density"), col = c("black", "red"), lty = c(1, 2))
# get the MAP value
density(sigmsq)$x[density(sigmsq)$y == max(density(sigmsq)$y)] -> v2
abline(v=v2)
# get actual MLE
av2 = sigmsq.seq[dinvchisq(sigmsq.seq, n-1, sy^2) == max(dinvchisq(sigmsq.seq, n-1, sy^2))]
abline(v = av2, col = "red", lty = 2)
```



```
# get absolute differences
abs(av1 - v1); abs(av2 - v2)
```

```
## [1] 0.001222942
```

```
## [1] 0.004788182
```

μ MC estimate flattens out so I went ahead and tried to look for an actual number by essentially looking for the mode. I eventually decided to do this for both estimators and both True and MC estimates. At the end I took the absolute difference and because μ has the smaller difference, I took that to be the one that is better estimate. Nevertheless, they both get relatively close.

```
####(d)
```

```
# get y's
ypredpost = c()
for(i in 1:B){
  ypredpost[i] = rnorm(1,mean=mu[i],sd = sqrt(sigmsq[i]))
}
```

recall from hw2 that

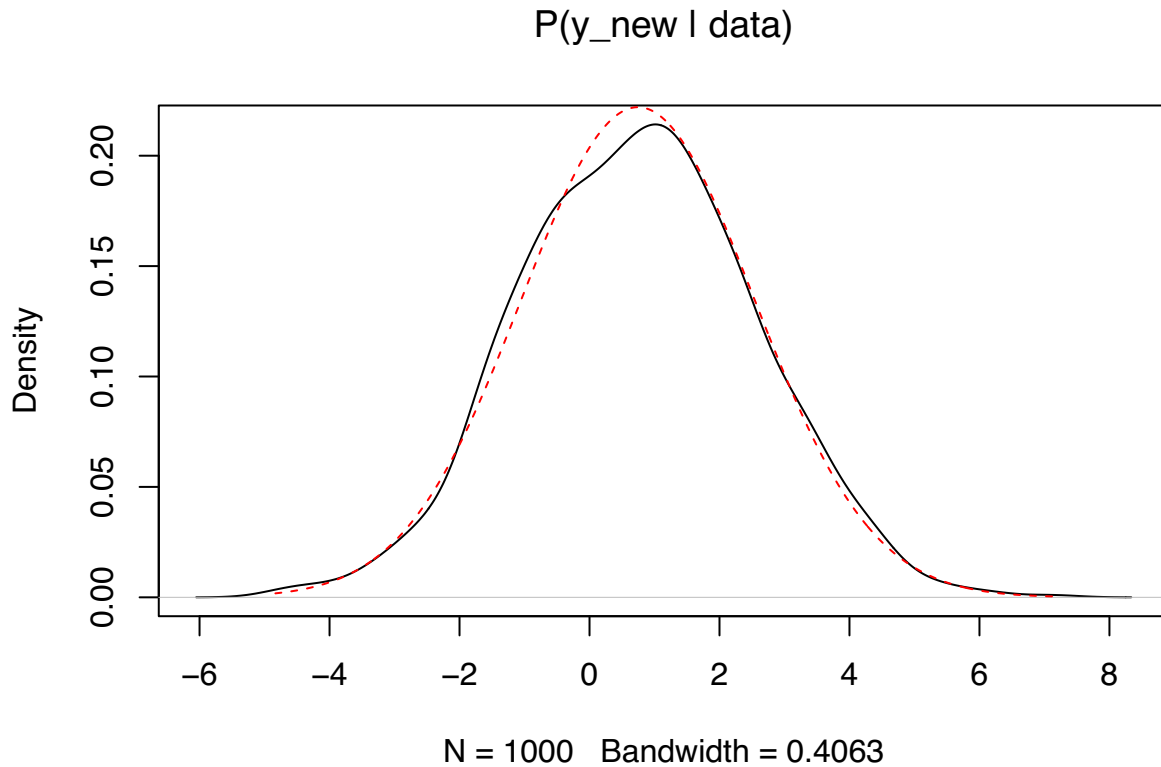
$$P(\mu, \sigma^2 | y_i) \propto (\sigma^2)^{-2} \prod_{i=1}^n \left[(\sigma^2)^{1/2} \exp \left(-\frac{1}{2} \left(\frac{y_i - \mu}{\sigma} \right)^2 \right) \right]$$

so that

$$\begin{aligned} P(\tilde{y} | y_i) &= \int_{R^+} \int_R N(\tilde{y} | \mu, \sigma^2) P(\mu, \sigma^2 | y_i) d\mu d\sigma^2 \\ &= t_{n-1}(\tilde{y}, s_y^2/n) \end{aligned}$$

which is the t-distribution

```
# plot density of MC samples
plot(density(ypredpost), main = expression(paste('P(y_new | data)', sep = "")))
# overlay "true" values
y.seq = seq(min(ypredpost), max(ypredpost), length.out = B)
lines(y.seq, dnorm(y.seq, mean(ypredpost), sd(ypredpost)), col = "red", lty = 2)
```



####(e)

The conditionals

$$\begin{aligned}\mu|\sigma^2, data &\sim N(\mu, \sigma^2/\kappa_n) \\ \sigma^2|\mu, data &\sim Sc.Inv - \chi^2(\nu'_n, \sigma_n^2)\end{aligned}$$

This is possible because we did it in class. I am going to write the algorithm in latex

Gibbs implies the following

1] $\sigma_b^2 \sim P(\sigma^2|\mu_{b-1}, data)$ 2] $\mu_b \sim P(\mu|\sigma_b^2, data)$ 3] Move back to 1 and restart

to code this, we need to use the “full posterior conditional distributions” for μ and σ^2 , the algorithm becomes

1] get the variance. (a.) $\frac{k_0(\mu_{b-1}-\mu_0)^2 + v_0\sigma_0^2 + (n-1)s_y^2 + n(\bar{y}-\mu_0)^2}{v'_n} \rightarrow \sigma_n^2$ (b.) $\sigma_b^2|\mu, data \sim Sc.Inv - \chi^2(\nu'_n, \sigma_n^2)$
 2] get the expected value $\mu_b \sim P(\mu|\sigma_b^2, data) \leftarrow$ feed this value back into (1a).

```
# I will pick initial values.
# It will only be simple.
```

```
k0 = 1
mu0 = 1
nu0 = 1
sig20 = 1
```

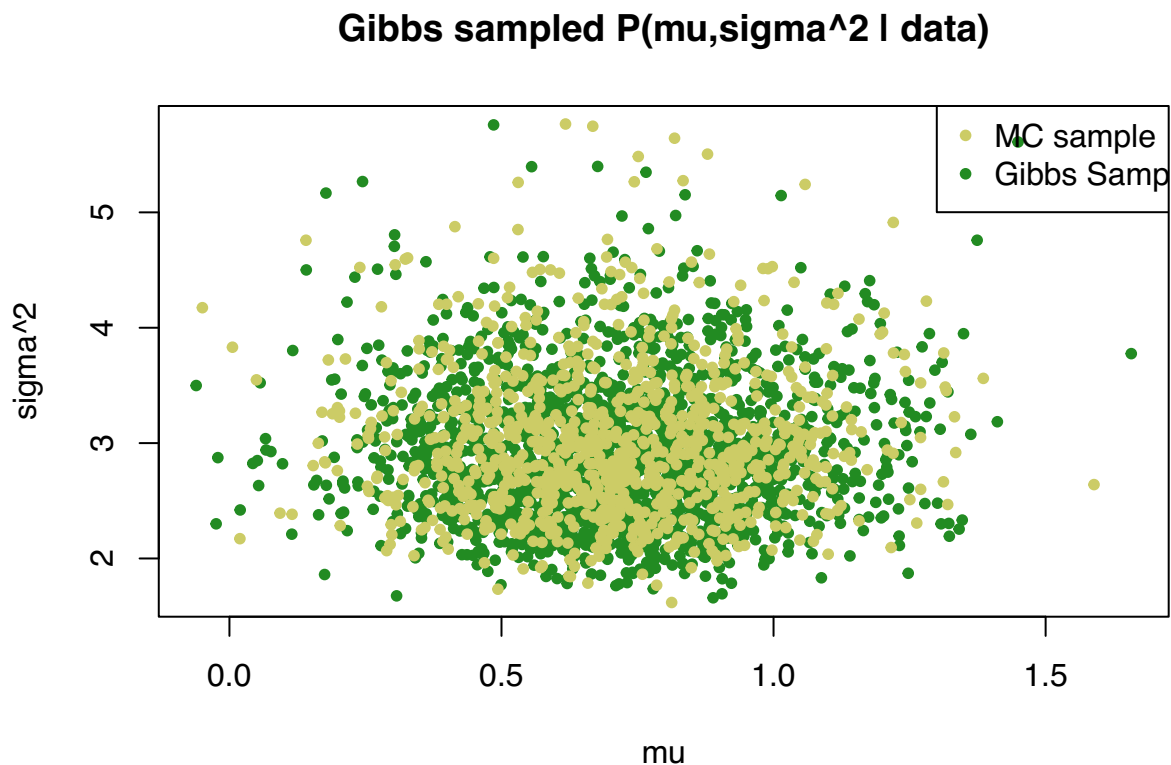
```

mu_n <- (k0/(k0 + n) * mu0) + (n/(k0 + n) *
ybar)
k_n <- k0 + n
nu_n.p <- nu0 + n + 1
#
## Posterior samples of mu and sigma2 initialize
mu.post.gibbs <- numeric()
sig2.post.gibbs <- numeric()
mu.post.gibbs[1] <- mean(y)
B <- 2000

### Running the Gibbs Sampler
for (b in 2:(B + 1)) {
sig2n.p <- (k0 * (mu.post.gibbs[(b - 1)] - mu0)^2 + nu0 * sig20 + (n - 1) * sy^2 + (n * (ybar - mu.post.
sig2.post.gibbs[(b)] <- rinvcchsq(1, df = nu_n.p, scale = sig2n.p) #sig2/ mu,data
mu.post.gibbs[(b)] <- rnorm(1, mean = mu_n, sd = sqrt(sig2.post.gibbs[(b)]/k_n)) #mu/sig2,data
}

plot(mu.post.gibbs,sig2.post.gibbs,
     main = "Gibbs sampled P(mu,sigma^2 | data)",
     xlab = "mu", ylab = "sigma^2",
     col = "forestgreen", pch = 20)
points(mu,sigmsq, pch = 20, col = rgb(0.8,0.8,0.4,1))
legend(1.3,6,c("MC sample","Gibbs Sample"), col = c(rgb(0.8,0.8,0.4,1),"forestgreen"), pch = 20)

```



Looks like the sampling method are almost identical. MC sampling was much easier to do; that is perhaps why we will be utilizing it in the next exercise.

Question 2

```
# artichoke data
y = c(1.8,1.6,1.4,1.6,1.4,1.5,1.2)

####(a)
# hybrid sampling system

beta = numeric()
alpha = numeric()
alpha[1] = 1
beta[1] <- rgamma(1,alpha[1]*n + 1, 0.8 + sum(y))
n = length(y)
s.tune = 3
accept = 0

#pull an alpha out
# make functions
Prior.alpha = function(x){
  dgamma(x,2,0.5)
}
Likeli = function(data,A,B){
  prod(dgamma(data,A,B))
}
#J.b = function(data,A){dlnorm(data,A,s.tune)}
J.b = function(x,A){dgamma(x,A*s.tune,s.tune)}
B=200000
##### Gibb + MH Loop #####
for(b in 2:(B+1)){
  #pull a beta out
  beta[b] <- rgamma(1,alpha[b-1]*n + 1, 0.8 + sum(y))

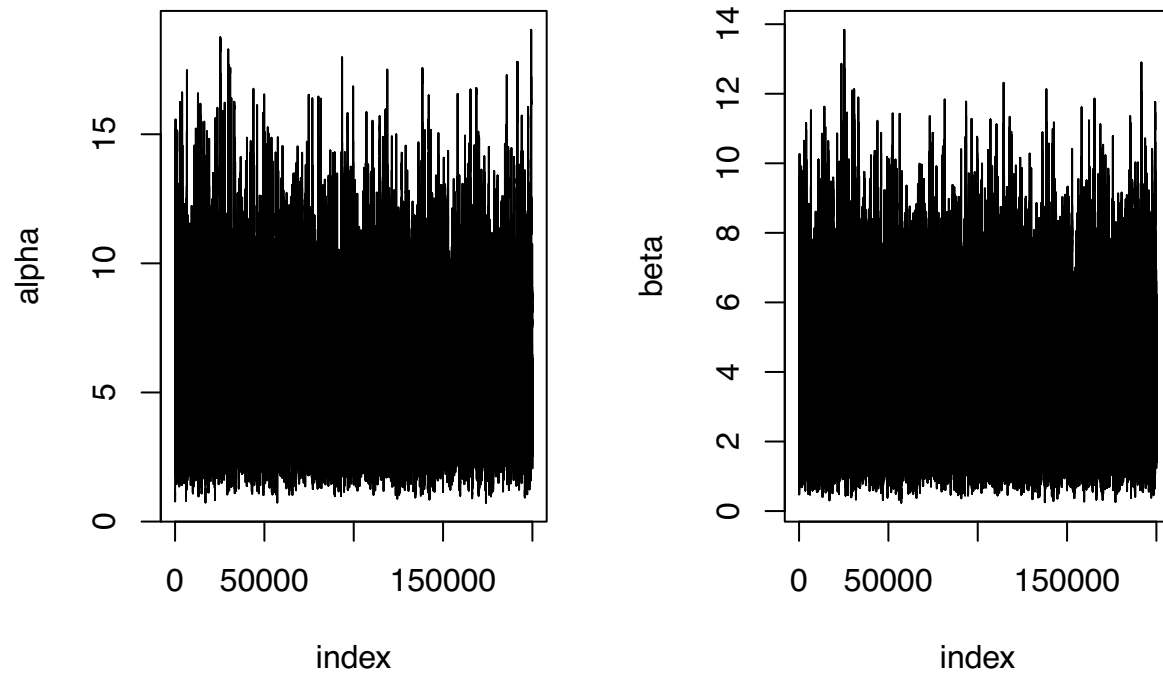
  # propose an alpha
  #rlnorm(1,alpha[b-1],s.tune) -> alpha.star
  rgamma(1,alpha[b-1]*s.tune,s.tune) -> alpha.star
  # create r
  num = Likeli(y,alpha.star,beta[b])*Prior.alpha(alpha.star)/J.b(alpha.star,alpha[b-1])
  denom = Likeli(y,alpha[b-1],beta[b])*Prior.alpha(alpha[b-1])/J.b(alpha[b-1],alpha.star)
  r = num/denom
  if(runif(1) < min(1,r)){
    alpha[b] = alpha.star
    accept = accept + 1
  } else {
    alpha[b] = alpha[b-1]
  }
}
accept/B

## [1] 0.55461
#####
```

I chose to propose from a Gamma distribution as the LogNorm gave me an acceptance rate of less than 0.5%.

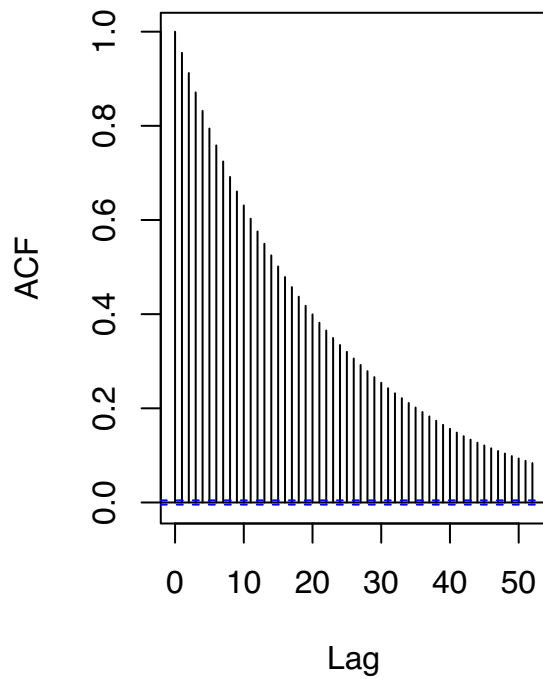
####(b.)

```
## trace plots
par(mfrow = c(1,2))
plot.ts(alpha,xlab = "index")
plot.ts(beta,xlab = "index")
```

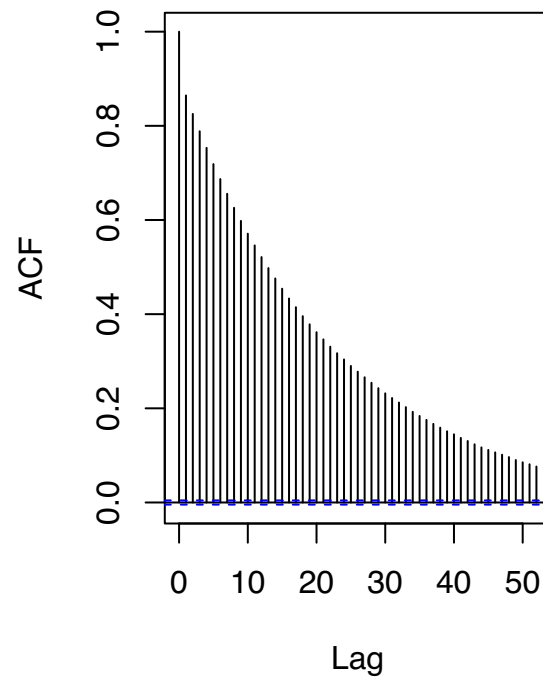


```
##autocorrelation plots
burnoff = 8000
par(mfrow = c(1,2))
acf(alpha[(burnoff+1):B])
acf(beta[(burnoff+1):B])
```

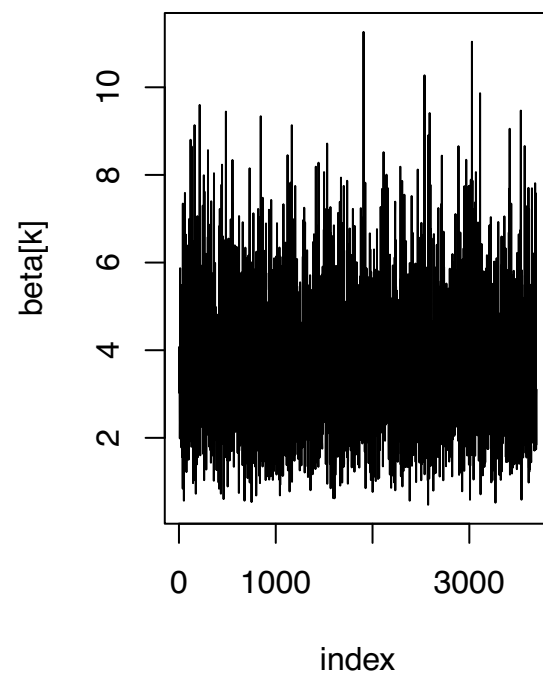
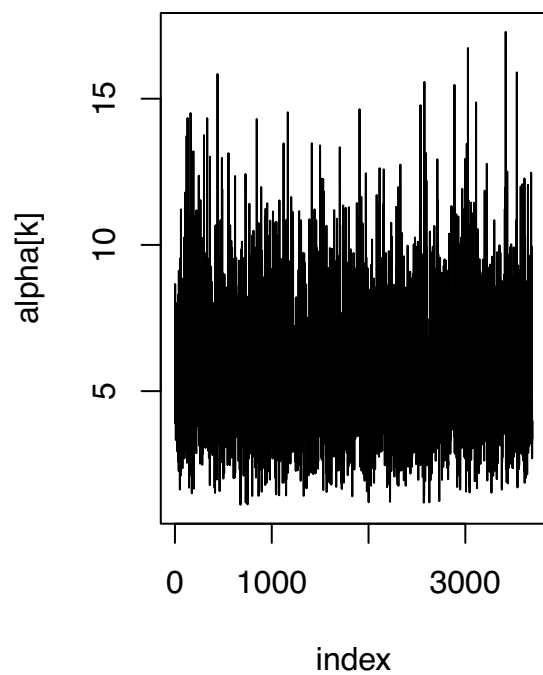
Series $\alpha[(\text{burnoff} + 1):B]$



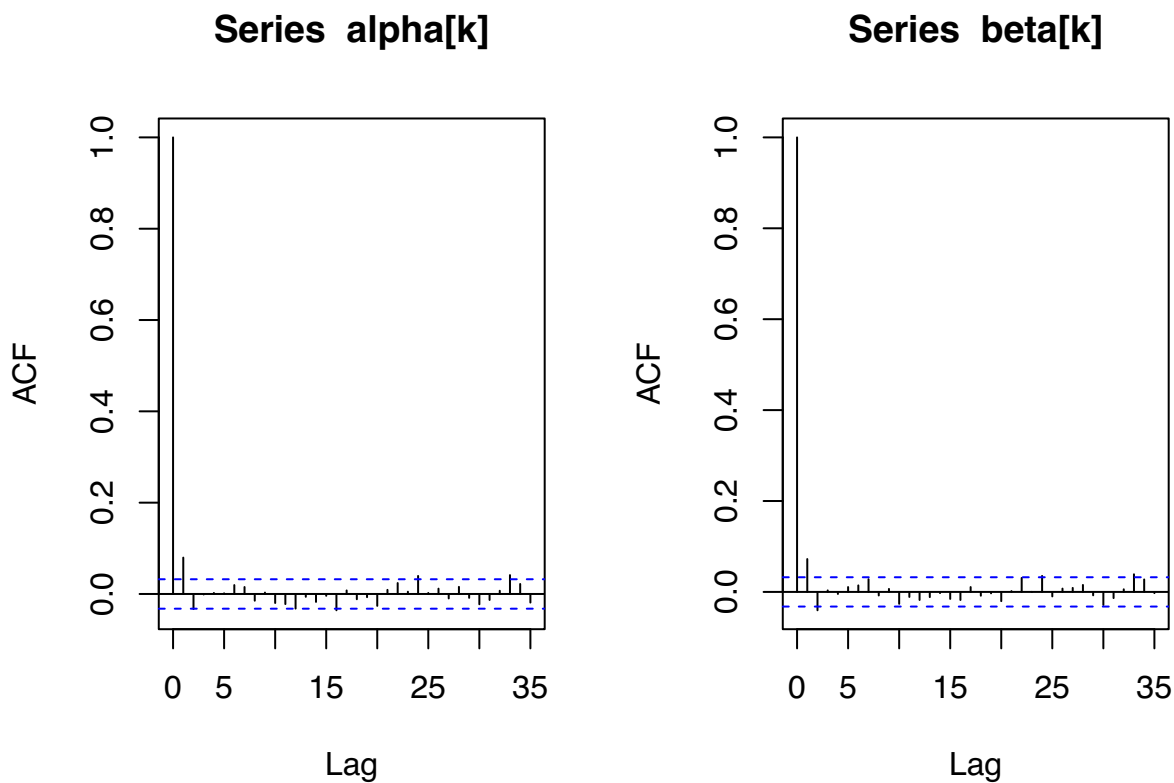
Series $\beta[(\text{burnoff} + 1):B]$



```
# grab every 52nd lag
k = 52*(1:B)
k = k[k < B & k > burnoff]
par(mfrow = c(1,2))
plot.ts(alpha[k],xlab = "index")
plot.ts(beta[k],xlab = "index")
```



```
acf(alpha[k])
acf(beta[k])
```



```
sprintf("I have obtained an acceptance rate of %f and I have burned off %d while also thinning by taking
```

```
## [1] "I have obtained an acceptance rate of 0.554610 and I have burned off 8000 while also thinning by
```

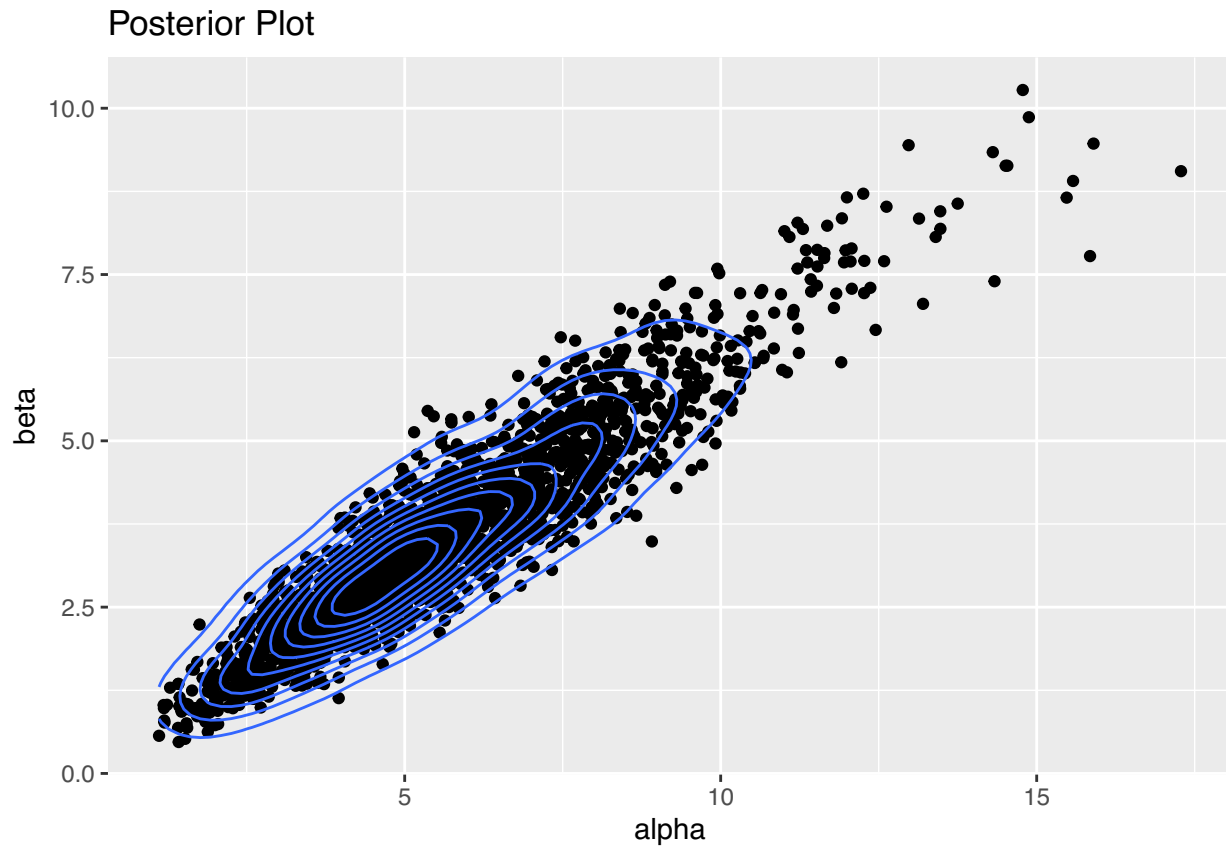
```
#####(d)
```

```
# sample 2000
```

```
k.i = sample(k, size = 2000)
```

```
post <- data.frame(cbind(alpha[k.i],beta[k.i]))
```

```
ggplot(post, aes(x = alpha[k.i], y =beta[k.i])) + geom_point() + geom_density2d() + ylab("beta") + xlab
```



```
alpha[k.i] -> alpha
beta[k.i] -> beta
al.MAP = density(alpha)$x[density(alpha)$y == max(density(alpha)$y)]
be.MAP = density(beta)$x[density(beta)$y == max(density(beta)$y)]
al.MAP
```

```
## [1] 4.875447
```

```
be.MAP
```

```
## [1] 3.112147
```

```
quantile(alpha,.025);quantile(alpha,.975)
```

```
##      2.5%
```

```
## 1.988586
```

```
##      97.5%
```

```
## 11.15699
```

```
quantile(beta,.025);quantile(beta,.975)
```

```
##      2.5%
```

```
## 1.169748
```

```
##      97.5%
```

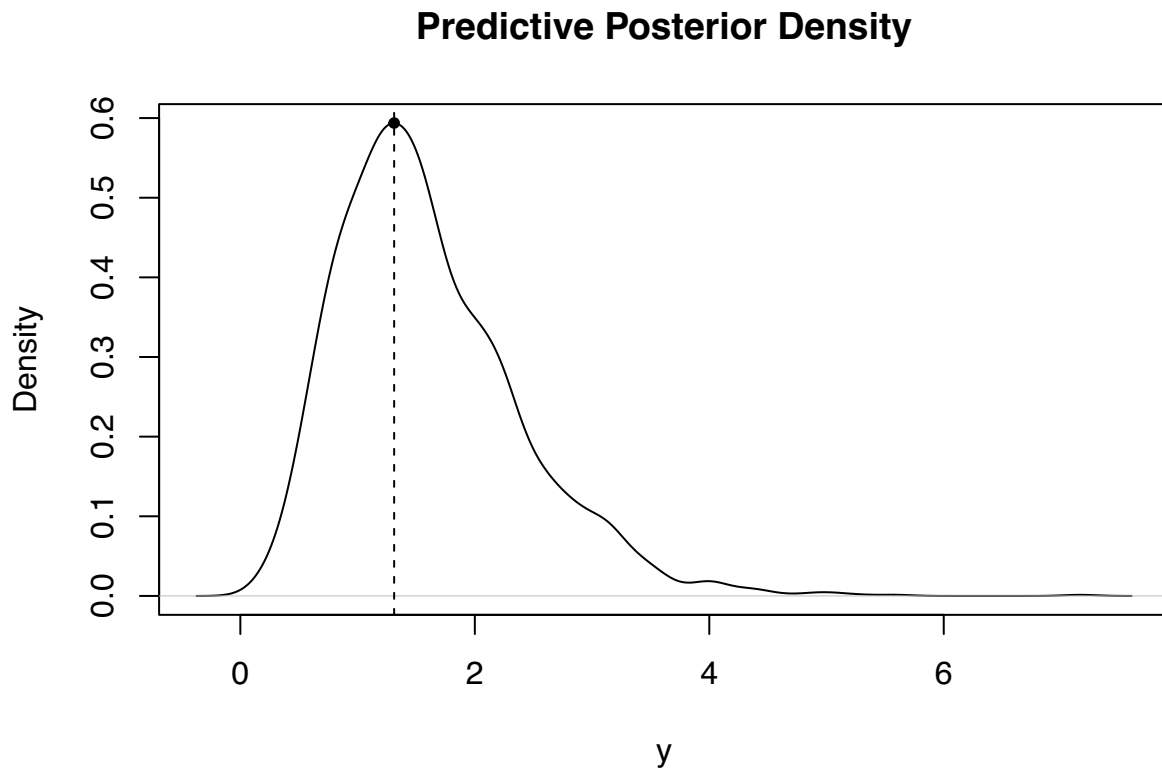
```
## 7.223074
```

	MAP	Cred. Int.
α	5.031591	(2.143953,10.78323)
β	3.113442	(1.278739,6.860119)

```

ynews = c()
for(i in 1:2000){
  ynews[i] = rgamma(1,alpha[i],beta[i])
}
ydens = density(ynews)
ynewMAP = ydens$x[ydens$y == max(ydens$y)]
plot(ydens,
     main = "Predictive Posterior Density",
     xlab = "y")
points(ynewMAP,ydens$y[ydens$y == max(ydens$y)],pch = 20)
abline(v = ynewMAP,lty = 2)

```



(e)

```

sprintf("I have found the MAP of Ynew to be %f with a 95-percent credible interval of (%f,%f)",ynewMAP,

```

```

## [1] "I have found the MAP of Ynew to be 1.311727 with a 95-percent credible interval of (0.480522,3.

```