

hw4

Michael Pena

2024-12-07

problem 13

```
# setting up data
bike <- read.csv("bike-data.csv", header = T)

# in class code with MH and gibbs sampling
# predef the beyes functions
Prior <- function(a,b){
  (a + b)^(-5/2)
}

LLH <- function(theta,a,b){
  sum(log(dbeta(theta,a,b)))
}

Proposal <- function(a,b){ #Jacobian
  1/(a*b)
}

rProposal <- function(n,mean,cov){
  rmvnorm(n,mean,cov)
}

# build a function just for this algorithm
MHGIBBs <- function(y,N,B,alpha0,beta0,S.tune = diag(2)){
  # initializations
  J = length(y)
  accept = 0
  alpha.post = beta.post = numeric()
  theta.post = matrix(0,J,B)
  theta0 <- numeric(length = J)

  #loop
  for(b in 1:B){
    # Gibbs Step for theta
    for(j in 1:J){
      shp1 = alpha0 + y[j]
      shp2 = beta0 + N[j] - y[j]
      theta0[j] = rbeta(1, shp1, shp2)
    }
    # Metro-Haste step for alpha and beta
    phi1 = rProposal(1, c(log(alpha0),log(beta0)), 1*S.tune)

    alpha1 = exp(phi1[1])
    beta1 = exp(phi1[2])

    r = exp(
      LLH(theta0,alpha1,beta1)
      + log(Prior(alpha1,beta1))
      + log(Proposal(alpha0,beta0))
      - LLH(theta0,alpha0,beta0)
      - log(Prior(alpha0,beta0))
      - log(Proposal(alpha1,beta1)))
  }
}
```

```

    ## accept check
    if(runif(1) < min(1,r)){
      alpha0 = alpha1
      beta0 = beta1
      accept = accept + 1
    }
    # drop off the samplings
    alpha.post[b] <- alpha0
    beta.post[b] <- beta0
    theta.post[,b] <- theta0
  }
  # tuning the covariance matrix
  S.tune <- matrix(0,2,2)
  S.tune[2,1] <- S.tune[1,2] <- cov(log(alpha.post),log(beta.post))
  S.tune[1,1] <- var(log(alpha.post))
  S.tune[2,2] <- var(log(beta.post))

  print(accept/B)
  # attributes in the function
  return(list("alpha" = alpha.post, "beta" = beta.post, "theta" = theta.post, "AR" = accept/B, "S" = S.tune))
}



# let's run our functions
y <- bike$Bicycles
N <- bike$Bicycles + bike$OtherVehicles
mean(y/N) # this is about .2, make alpha0 = 2, beta0 = 8

## [1] 0.1961412
MHGIBBs(y,N,5000,2,8)$S -> S1 # run 1 time to get tuning matrix


## [1] 0.124
MHGIBBs(y,N,10000,2,8,S1) -> obj1

## [1] 0.5019
# visualizations
par(mfrow = c(2,2))
plot.ts(obj1$alpha)
plot.ts(obj1$beta)
acf(obj1$alpha)
acf(obj1$beta)


```



Series obj1\$alpha




Series obj1\$beta




```
# take out all but the 10th lag
k = 10*(1:10000)
k = k[k <= 10000]
```



```
# visualizations
par(mfrow = c(2,2))
plot.ts(obj1$alpha[k])
plot.ts(obj1$beta[k])
acf(obj1$alpha[k])
acf(obj1$beta[k])
```



Series obj1\$alpha[k]





Series obj1\$beta[k]





```
par(mfrow = c(2,2))
for(d in 1:10){
  plot(density(obj1$theta[d,k]),main = paste("Density of theta_",d, "with raw proportion",d))
  abline(v = (y/N)[d])
}
```


Density of theta_ 1 with raw proportion Density of theta_ 2 with raw proportion




Density of theta_ 3 with raw proportion Density of theta_ 4 with raw proportion



Density of theta_ 5 with raw proportion Density of theta_ 6 with raw proportion




N = 1000 Bandwidth = 0.007276




N = 1000 Bandwidth = 0.01026

Density of theta_ 7 with raw proportion Density of theta_ 8 with raw proportion




N = 1000 Bandwidth = 0.007683




N = 1000 Bandwidth = 0.006378

Density of theta_ 9 with raw proportion Density of theta_ 10 with raw proportion




N = 1000 Bandwidth = 0.004155



N = 1000 Bandwidth = 0.009863

```
plot(density(as.vector(obj1$theta[,k])))
```

density(x = as.vector(obj1\$theta[, k]))



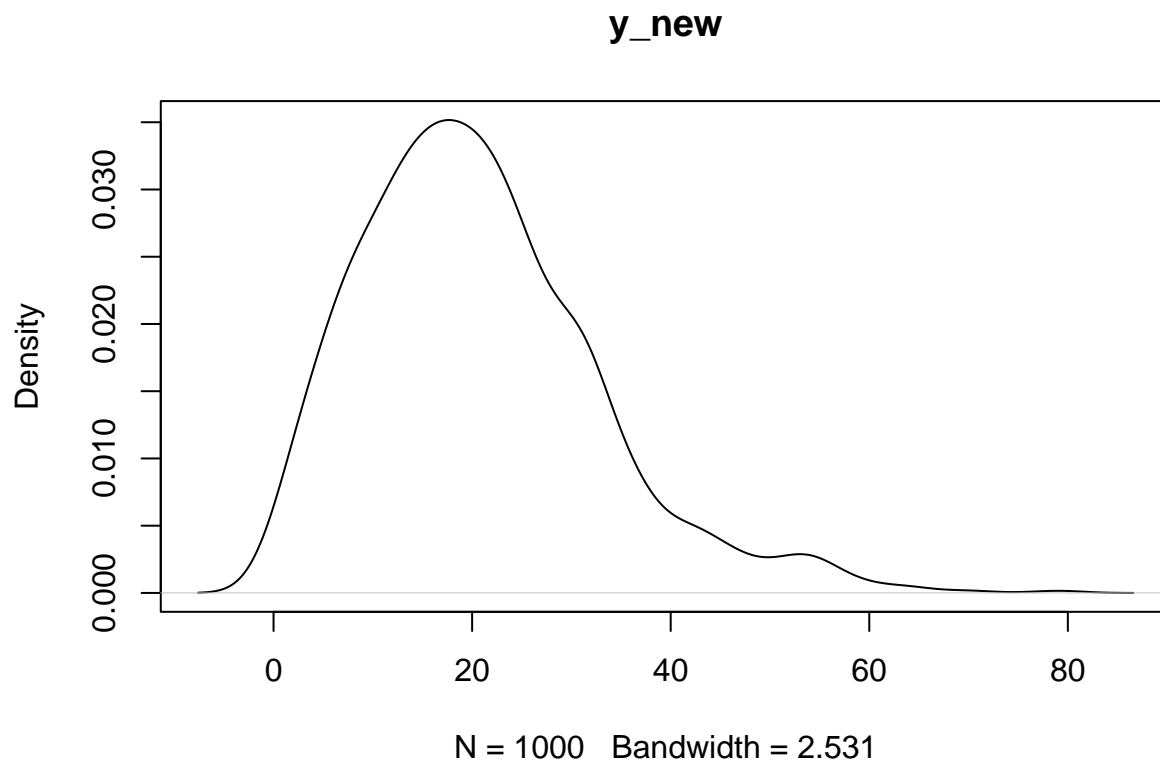
N = 10000 Bandwidth = 0.01007

```
quantile(obj1$theta[,k], c(0.025,.975))
```

```
##      2.5%      97.5%  
## 0.08293409 0.45729190
```

```
# input alphas and betas into a beta distribution  
rbeta(1000,obj1$alpha[k],obj1$beta[k]) -> newtheta  
rbinom(1000,100,newtheta) -> newy
```

```
plot(density(newy), main = "y_new")
```



```
quantile(newy,c(.025,.975))
```

```
## 2.5% 97.5%
##    2    52
```

```
# checking analytically if this makes sense
```

```
CI = matrix(0, ncol = 3, nrow = 10)
```

```
for(j in 1:10){
```

```
  CI[j,] = quantile(obj1$theta[j,k], probs = c(0.025,0.5, 0.975))
```

```
}
```


```
plot(y/N, y/N, type = "l")
```

```
points(y/N, CI[,2], pch = 19, col = 3)
```

```
for(j in 1:10){
```

```
  points(c(y[j]/N[j],y[j]/N[j]), c(CI[j,1],CI[j,3]), type = "l", col = 4)
```

```
}
```

This distribution is pretty reasonable according to this graph

Additional Problem

```
# setup the data
schools <- read.csv(file = "schools.csv", header = T)
schools <- list("J" = 8,
               "y" = schools$estimate,
               "sigma" = schools$sd)
```

```
# run the STAN and fit the data
# schools_fit <- stan(file="schools.stan",
#                   data = schools,
#                   iter = 1000,
#                   chains = 4)
fit1 <- stan(
  file = "schools.stan", # Stan program
  data = schools,        # named list of data
  chains = 4,           # number of Markov chains
  warmup = 1000,        # number of warmup iterations per chain
  iter = 20000,         # total number of iterations per chain
  cores = 2,            # number of cores
  refresh = 1000,       # show progress every 'refresh' iterations
  thin = 10             # number of thinning
)
```

```
## Trying to compile a simple C file
```

```
## Running /usr/lib/R/bin/R CMD SHLIB foo.c
## using C compiler: 'gcc (Ubuntu 11.4.0-1ubuntu1~22.04) 11.4.0'
## gcc -I"/usr/share/R/include" -DNDEBUG -I"/home/cern/R/x86_64-pc-linux-gnu-library/4.4/Rcpp/include"
## In file included from /home/cern/R/x86_64-pc-linux-gnu-library/4.4/RcppEigen/include/Eigen/Core:19,
## from /home/cern/R/x86_64-pc-linux-gnu-library/4.4/RcppEigen/include/Eigen/Dense:1,
```

```

##          from /home/cern/R/x86_64-pc-linux-gnu-library/4.4/StanHeaders/include/stan/math/pri
##          from <command-line>:
## /home/cern/R/x86_64-pc-linux-gnu-library/4.4/RcppEigen/include/Eigen/src/Core/util/Macros.h:679:10:
## 679 | #include <cmath>
##     |         ^~~~~~
## compilation terminated.
## make: *** [/usr/lib/R/etc/Makeconf:195: foo.o] Error 1

## Warning: There were 7 divergent transitions after warmup. See
## https://mc-stan.org/misc/warnings.html#divergent-transitions-after-warmup
## to find out why this is a problem and how to eliminate them.

## Warning: Examine the pairs() plot to diagnose sampling problems

print (fit1)

## Inference for Stan model: anon_model.
## 4 chains, each with iter=20000; warmup=1000; thin=10;
## post-warmup draws per chain=1900, total post-warmup draws=7600.
##
##          mean se_mean  sd  2.5%  25%  50%  75% 97.5% n_eff Rhat
## mu          7.87    0.06 5.19  -2.12  4.57  7.86 11.05 18.36 6446  1
## tau          6.55    0.06 5.52   0.25  2.55  5.27  9.09 20.59 7352  1
## eta[1]       0.39    0.01 0.94  -1.54 -0.22  0.41  1.03  2.18 7483  1
## eta[2]       0.00    0.01 0.87  -1.71 -0.57 -0.01  0.58  1.72 7613  1
## eta[3]      -0.19    0.01 0.93  -1.98 -0.81 -0.20  0.45  1.64 7487  1
## eta[4]      -0.03    0.01 0.89  -1.83 -0.62 -0.04  0.56  1.73 7508  1
## eta[5]      -0.35    0.01 0.88  -2.02 -0.94 -0.38  0.22  1.46 7013  1
## eta[6]      -0.20    0.01 0.90  -2.00 -0.79 -0.20  0.38  1.59 7399  1
## eta[7]       0.35    0.01 0.88  -1.43 -0.20  0.37  0.92  2.10 7696  1
## eta[8]       0.06    0.01 0.94  -1.79 -0.55  0.07  0.70  1.94 7502  1
## theta[1]    11.22    0.10 8.24  -2.57  5.93 10.15 15.55 30.71 7420  1
## theta[2]     7.84    0.07 6.23  -4.62  3.94  7.79 11.72 20.41 6948  1
## theta[3]     6.18    0.09 7.76 -11.29  2.13  6.69 10.91 20.54 7745  1
## theta[4]     7.57    0.08 6.52  -5.92  3.62  7.61 11.55 20.63 7530  1
## theta[5]     5.13    0.07 6.33  -8.62  1.28  5.50  9.36 16.73 7417  1
## theta[6]     6.16    0.08 6.66  -8.26  2.29  6.55 10.45 18.44 7515  1
## theta[7]    10.66    0.08 6.81  -1.07  5.99  9.98 14.64 26.40 7227  1
## theta[8]     8.41    0.10 8.03  -7.87  3.88  8.25 12.72 25.90 7079  1
## lp__        -4.89    0.03 2.65 -10.69 -6.46 -4.64 -2.99 -0.43 7622  1
##
## Samples were drawn using NUTS(diag_e) at Sat Dec  7 15:57:41 2024.
## For each parameter, n_eff is a crude measure of effective sample size,
## and Rhat is the potential scale reduction factor on split chains (at
## convergence, Rhat=1).

plot (fit1)

## 'pars' not specified. Showing first 10 parameters by default.
## ci_level: 0.8 (80% intervals)
## outer_level: 0.95 (95% intervals)

```

