

Math 538: Bayesian Statistics

Lecture Slides 6: Multiple Parameter Models with STAN

Department of Mathematics, California State University, Fullerton



- 1 STAN: A Brief Introduction
- 2 Bioassay Example Using rstan
- 3 Bioassay Example Using brms

STAN: A Brief Introduction

What is STAN?



- STAN is a probabilistic programming language that implements Bayesian approaches to statistical modeling
- Written in C++
- High-performance
- Interfaces with: R, Julia, Python, shell, Matlab, and STATA
- See <http://mc-stan.org/> for more information
- Many presentation slides from StanConn2023 available at <https://github.com/stan-dev/stancon2023>

Installing STAN through R (RSTAN)

For installing RSTAN on your operating system:

- For Windows, Mac, or Linux Users:

<https://github.com/stan-dev/rstan/wiki/#installing-rstan>

Note that you will want to:

- 1 Update your operating system
- 2 Using R version 4.2.0 or later is strongly recommended
- 3 RStudio version 1.4x or later
- 4 You will need an updated C++ compiler on your computer (I use XCode)

Configuring C Toolchain

- Mac <https://github.com/stan-dev/rstan/wiki/Configuring-C---Toolchain-for-Mac>
- Windows <https://github.com/stan-dev/rstan/wiki/Configuring-C---Toolchain-for-Windows>
- Linux <https://github.com/stan-dev/rstan/wiki/Configuring-C-Toolchain-for-Linux>

```
# install macrotools
install.packages("remotes")
remotes::install_github("coatless-mac/macrtools")

# installs: Xcode CLI, gfortran, R Development
# binaries will ask for your computer password
macrtools::macos_rtools_install()
```

Congratulations!

Xcode CLI, Gfortran, and R developer binaries have been installed successfully.

Optimize and then install rstan

```

# Once the toolchain is installed you can then
# enable some compiler optimizations to improve
# the estimation speed of the model:
dotR <- file.path(Sys.getenv("HOME"), ".R")
if (!file.exists(dotR)) dir.create(dotR)
M <- file.path(dotR, "Makevars")
if (!file.exists(M)) file.create(M)
arch <- ifelse(R.version$arch == "aarch64", "arm64",
               "x86_64")
cat(paste("\nCXX14FLAGS += -O3 -mtune=native -arch",
          arch, "-ftemplate-depth-256"), file = M, sep = "\n",
    append = FALSE)

#-----
#-----Installing rstan -----
remove.packages("rstan")
if (file.exists(".RData")) file.remove(".RData")
Sys.setenv(MAKEFLAGS = "-j4") # four cores used
install.packages(c("Rcpp", "RcppEigen", "RcppParallel",
                  "StanHeaders"), type = "source")
install.packages("rstan", type = "source")

```

Run this example to make sure your rstan works

```
example(stan_model, package = "rstan", run.dontrun = TRUE)
```

```
tn_md> stancode <- 'data {real y_mean;} parameters {real y;}
  model {y ~ normal(y_mean,1);}'
```

```
stn_md> mod <- stan_model(model_code = stancode, verbose = TRUE)
```

TRANSLATING MODEL '' FROM Stan CODE TO C++ CODE NOW.

```
stn_md> fit <- sampling(mod, data = list(y_mean = 0))
```

SAMPLING FOR MODEL 'anon_model' NOW (CHAIN 1).

Chain 1:

Chain 1: Gradient evaluation took 3e-06 seconds

Chain 1: 1000 transitions using 10 leapfrog steps per transition
would take 0.03 seconds.

Chain 1: Adjust your expectations accordingly!

Chain 1:

Chain 1:

Chain 1: Iteration: 1 / 2000 [0%] (Warmup)

Now you can load

```
library("rstan")
```

*# If you are running on a local multi-core
processor do*

```
options(mc.cores = parallel::detectCores())  
rstan_options(auto_write = TRUE)
```

Bioassay Example Using rstan

Bioassay Example Using rstan

- Recall the Bioassay Problem p.74 of BDA

Dose (logscale)	# animals	# deaths
x_i	n_i	y_i
-0.86	5	0
-0.30	5	1
-0.05	5	3
0.73	5	5

- A standard model for these data are given by:

$$\begin{aligned}
 y_i | \theta_i &\sim \text{Bin}(n_i, \theta_i) \text{ for } i = 1, \dots, K \\
 \text{with } \text{logit}(\theta_i) &= \alpha + \beta x_i \\
 p(\alpha, \beta) &\propto 1
 \end{aligned}$$

In RStudio open a stan file and write:

```
data {
  int<lower=0> J;          // number of batches
  int<lower=0> n[J]; // number in each batch
  int<lower=0> y[J]; // number of deaths in each batch
  vector[J] x;           // log-dose of each batch
}
parameters {
  real alpha;           // intercept
  real beta;            // coefficient of dose effect
}
transformed parameters {
  vector[J] theta = inv_logit(alpha+beta*x); // death rate by dosage
}
model {
  y ~ binomial(n,theta); // log-likelihood
  //Note a uniform prior is assumed when prior not specified
}
```

```
#create data
bioassay_dat <- list(
  J = 4,
  n = c(5, 5, 5, 5),
  y = c(0, 1, 3, 5),
  x = c(-0.86, -0.30, -0.05, 0.73))

fit1 <- stan(
  file = "bioassay.stan", # Stan program
  data = bioassay_dat,    # named list of data
  chains = 4,             # number of Markov chains
  warmup = 1000,          # number of warmup iterations per chain
  iter = 20000,           # total number of iterations per chain
  cores = 2,              # number of cores
  refresh = 1000,         # show progress every 'refresh' iterations
  thin = 10               # number of thinning
)
```

Chain 4: Iteration: 20000 / 20000 [100%] (Sampling)


Chain 4:

Chain 4: Elapsed Time: 0.072 seconds (Warm-up)


Chain 4: 1.184 seconds (Sampling)

Some useful automatic plotting functions in rstan


```
stan_trace(fit1, pars = c("alpha", "beta"), inc_warmup = TRUE,  
           nrow = 2)
```




```
stan_ac(fit1, pars = c("alpha", "beta"), inc_warmup = FALSE,  
        nrow = 2, separate_chains = TRUE, lags = 25)
```



```
pairs(fit1, pars = c("alpha", "beta"), las = 1)
```




```
fit_ss <- extract(fit1)
names(fit_ss)
# [1] 'alpha' 'beta' 'theta' 'lp__'
biv_kde <- MASS::kde2d(fit_ss$alpha, fit_ss$beta)
plot(fit_ss$alpha, fit_ss$beta, pch = ".", xlab = "alpha",
     ylab = "beta")
contour(biv_kde, col = 2, add = T)
```



```
print(fit1, pars = c("alpha", "beta", "theta"), probs = c(0.025,
  0.5, 0.975))
```

Inference for Stan model: anon_model.

4 chains, each with iter=20000; warmup=1000; thin=10;


post-warmup draws per chain=1900, total post-warmup draws=7600.

	mean	se_mean	sd	2.5%	50%	97.5%	n_eff	Rhat
alpha	1.32	0.01	1.10	-0.54	1.23	3.78	7448	1
beta	11.66	0.07	5.87	3.45	10.64	25.51	7139	1
theta[1]	0.01	0.00	0.02	0.00	0.00	0.07	7443	1
theta[2]	0.15	0.00	0.13	0.01	0.12	0.47	6931	1
theta[3]	0.65	0.00	0.18	0.29	0.66	0.94	7569	1
theta[4]	0.99	0.00	0.03	0.92	1.00	1.00	7468	1

Samples were drawn using NUTS(diag_e) at Sun Oct 29 22:50:20 2023.

For each parameter, n_eff is a crude measure of effective sample size, and Rhat is the potential scale reduction factor on split chains (at convergence, Rhat=1).

```
plot(fit1)
```




Inference for LD50

Recall: LD50: $E\left(\frac{y_i}{n_i}\right) = \text{logit}^{-1}(\alpha + \beta x_i) = 0.5 \Rightarrow$ occurs when $x_i = -\alpha/\beta$

Posterior of LD50

```
LD50 = -fit_ss$alpha/fit_ss$beta
quantile(LD50, probs = c(0.25, 0.5, 0.975))
plot(density(LD50), type = "l", main = "LD50")
```

25%	50%	97.5%
-0.16217623	-0.11249260	0.09887022



Bioassay Example Using brms

brms: Bayesian Regression Models Using Stan

The brms package provides an interface to fit Bayesian generalized (non)linear multivariate multilevel models using Stan.

```
install.packages("brms")  
library("brms")
```

```
bioassay_dat = data.frame(n = c(5, 5, 5, 5), y = c(0,  
  1, 3, 5), x = c(-0.86, -0.3, -0.05, 0.73))
```

```
fit2 <- brm(y | trials(n) ~ x, data = bioassay_dat,  
  family = binomial())
```

```
Chain 4: Elapsed Time: 0.091 seconds (Warm-up)
```

```
Chain 4:           0.08 seconds (Sampling)
```

```
Chain 4:           0.171 seconds (Total)
```

`summary(fit2)``Family: binomial``Links: mu = logit``Formula: y | trials(n) ~ x``Data: bioassay_dat (Number of observations: 4)``Draws: 4 chains, each with iter = 2000; warmup = 1000; thin = 1;
total post-warmup draws = 4000``Population-Level Effects:`

	Estimate	Est.Error	l-95% CI	u-95% CI	Rhat	Bulk_ESS
Intercept	1.35	1.09	-0.48	3.82	1.00	1606
x	11.75	5.87	3.54	25.82	1.00	1414

`Tail_ESS`

Intercept	1868
x	1378

Draws were sampled using `sampling(NUTS)`. For each parameter, `Bulk_ESS` and `Tail_ESS` are effective sample size measures, and `Rhat` is the potential scale reduction factor on split chains (at convergence, `Rhat` = 1).

```
plot(fit2)
```

