


Quiz 2 Takehome

Michael Pena

2024-11-26

Question 1

```
boxplot(hear)
```



(a)

```
summary(hear)
```

| | | | | |
|----|---------------|---------------|---------------|---------------|
| ## | List1 | List2 | List3 | List4 |
| ## | Min. :18.00 | Min. :16.00 | Min. :14.00 | Min. :14.00 |
| ## | 1st Qu.:28.00 | 1st Qu.:21.50 | 1st Qu.:19.50 | 1st Qu.:19.50 |
| ## | Median :32.00 | Median :30.00 | Median :25.00 | Median :25.00 |
| ## | Mean :32.75 | Mean :29.67 | Mean :25.25 | Mean :25.58 |
| ## | 3rd Qu.:36.50 | 3rd Qu.:36.00 | 3rd Qu.:32.00 | 3rd Qu.:30.50 |
| ## | Max. :48.00 | Max. :44.00 | Max. :38.00 | Max. :42.00 |

Right off the bat, we notice that there is a list 1 and list 2 have completely different spreads. List 1 is also noticeably higher and less spread than all the other lists. List 2 has the most spread and it looks as though list 3 and 4 are very similar. we could possibly start making our assumptions now about assuming the effect that being in certain list implies on the score.

(b)

$$\begin{aligned} P(\theta, \mu, \sigma^2 | \mathbf{y}) \\ = P(\mu) P(\sigma^2) \prod_{j=1}^J P(\theta_j | \mu, \sigma^2) \prod_{j=1}^J \prod_{i=1}^{n_j} P(y_{ij} | \theta_j, \sigma^2) \end{aligned}$$

(c) we care about θ_j so we extract $P(\theta_j|\mu, \sigma^2) \prod_{i=1}^{n_j} P(y_{ij}|\theta_j, \sigma^2)$

$$\begin{aligned}
& P(\theta_j|\mu, \sigma^2) \prod_i^{n_j} P(y_{ij}|\theta_j, \sigma^2) \\
& \propto \exp\left[\frac{-1}{2\sigma^2}((\theta_j - \mu)^2 + \sum_i^{n_j} (y_{ij} - \theta_j)^2)\right] \\
& = \exp\left[\frac{-1}{2\sigma^2}(\theta_j^2 - 2\mu\theta_j + \mu^2 + \sum_i^{n_j} (y_{ij}^2 - 2\theta_j y_{ij} + \theta_j^2))\right] \\
& \propto \exp\left[\frac{-1}{2\sigma^2}(\theta_j^2 - 2\mu\theta_j + n_j\theta_j^2 - 2\theta_j \sum_i^{n_j} y_{ij})\right] \\
& = \exp\left[\frac{-(n_j + 1)}{2\sigma^2}(\theta_j^2 - 2\theta_j \frac{\mu + \sum_i^{n_j} y_{ij}}{n_j + 1} + (\frac{\mu + \sum_i^{n_j} y_{ij}}{n_j + 1})^2 - (\frac{\mu + \sum_i^{n_j} y_{ij}}{n_j + 1})^2)\right] \\
& \propto \exp\left[\frac{-(n_j + 1)}{2\sigma^2}(\theta_j^2 - 2\theta_j \frac{\mu + \sum_i^{n_j} y_{ij}}{n_j + 1} + (\frac{\mu + \sum_i^{n_j} y_{ij}}{n_j + 1})^2)\right] \\
& = \exp\left[\frac{-(n_j + 1)}{2\sigma^2}(\theta_j - \frac{\mu + \sum_i^{n_j} y_{ij}}{n_j + 1})^2\right]
\end{aligned}$$

This gives us $N(\frac{\mu + \sum_i^{n_j} y_{ij}}{n_j + 1}, \frac{\sigma^2}{1 + n_j})$

Thus we have $P(\theta_j|\theta_{-j}, \mu, \sigma^2, Y_{ij}) \sim N(\frac{\mu + \sum_i^{n_j} y_{ij}}{25}, \frac{\sigma^2}{25})$

(d) finding the $\mu|\theta_j, \mu, \sigma^2, \mathbf{Y}$

$$\begin{aligned}
& P(\mu) \prod_j^4 P(\theta_j | \mu, \sigma^2) \\
& \propto \exp \left[-(\mu - 30)^2 - \frac{1}{2\sigma^2} \sum_i^4 (\theta_j - \mu)^2 \right] \\
& = \exp \left[- \left[\mu^2 - 60\mu + 900 + \frac{1}{2\sigma^2} \sum_j^4 \theta_j^2 - \frac{2\mu}{2\sigma^2} \sum_j^4 \theta_j + \frac{24}{2\sigma^2} \mu^2 \right] \right] \\
& \propto \exp \left[- \left[\mu^2 - 60\mu - \frac{2\mu}{2\sigma^2} \sum_j^4 \theta_j + \frac{24}{2\sigma^2} \mu^2 \right] \right] \\
& = \exp \left[- \left[\mu^2 \left(1 + \frac{12}{\sigma^2}\right) - \mu \left(60 + \frac{1}{\sigma^2} \sum_j^4 \theta_j\right) \right] \right] \\
& = \exp \left[- \frac{1}{1 + \frac{12}{\sigma^2}} \left[\mu^2 - \mu \left(\frac{60 + \frac{1}{\sigma^2} \sum_j^4 \theta_j}{1 + \frac{12}{\sigma^2}} \right) \right] \right]
\end{aligned}$$

note:

$$\begin{aligned}
& - \frac{1}{1 + \frac{12}{\sigma^2}} = \frac{-\sigma^2}{2(\frac{\sigma^2}{2} + 6)} \\
& \frac{60 + \frac{1}{\sigma^2} \sum_j^4 \theta_j}{1 + \frac{12}{\sigma^2}} = \frac{60\sigma^2 + \sum_j^4 \theta_j}{\sigma^2 + 12} \\
& = \exp \left[\frac{-\sigma^2}{2(\frac{\sigma^2}{2} + 6)} \left[\mu^2 - \mu \left(\frac{60\sigma^2 + \sum_j^4 \theta_j}{\sigma^2 + 12} \right) + \left[\frac{1}{2} \left(\frac{60\sigma^2 + \sum_j^4 \theta_j}{\sigma^2 + 12} \right)^2 - \left[\frac{1}{2} \left(\frac{60\sigma^2 + \sum_j^4 \theta_j}{\sigma^2 + 12} \right)^2 \right] \right] \right] \right] \\
& \propto \exp \left[\frac{\sigma^2}{2(\frac{\sigma^2}{2} + 6)} \left[\mu - \frac{1}{2} \left(\frac{60\sigma^2 + \sum_j^4 \theta_j}{\sigma^2 + 12} \right) \right]^2 \right] \\
& \Rightarrow \mu | \theta, \sigma^2, Y \sim \text{Norm} \left(\frac{1}{2} \left(\frac{60\sigma^2 + \sum_j^4 \theta_j}{\sigma^2 + 12} \right), \frac{\sigma^2}{\frac{\sigma^2}{2} + 6} \right)
\end{aligned}$$

finding the distribution of $\sigma^2 | \mu, \theta, \mathbf{Y}$

$$\begin{aligned}
& P(\sigma^2) \prod_j^4 P(\theta_j | \mu, \sigma^2) \prod_j^4 \prod_i^{24} P(y_{ij} | \theta_j, \sigma^2) \\
& = 100 \left(\frac{1}{\sigma^2} \right)^3 e^{\frac{-10}{\sigma^2}} \left[\prod_j^4 \prod_i^{24} \frac{1}{\sqrt{2\pi\sigma^2}} e^{\frac{-1}{2\sigma^2} (y_{ij} - \theta_j)^2} \right] \left[\prod_i^{24} \frac{1}{\sqrt{2\pi\sigma^2}} e^{\frac{-1}{2\sigma^2} (\theta_j - \mu)^2} \right] \\
& \propto \left(\frac{1}{\sigma^2} \right)^{24(5)+3} \exp \left[\frac{-1}{\sigma^2} \left(10 + \frac{1}{2} \sum_j^4 \sum_i^{24} (y_{ij} - \theta_j)^2 + \frac{1}{2} \sum_i^{24} (\theta_j - \mu)^2 \right) \right] \\
& \propto \text{InvGamma} \left(122, 10 + \frac{1}{2} \sum_j^4 \sum_i^{24} (y_{ij} - \theta_j)^2 + \frac{1}{2} \sum_i^{24} (\theta_j - \mu)^2 \right)
\end{aligned}$$

```

set.seed(538)
# making the MCMC algorithm
GIBBS <- function(y, B = 5000){
  # establish initials

```

```

nrow(y) -> nj
ncol(y) -> J
theta.post <- matrix(0,B,J)
mu.post <- sigsq.post <- numeric(length = B)
theta.post[1,] = as.vector(colMeans(y))
mu.post[1] = mean(theta.post[1,])
sigsq.post[1] = var(as.vector(y))

# MCMC loop
for(i in 2:B){
  # set/reset
  sigsq0 <- sigsq.post[i-1]
  mu0 <- mu.post[i-1]
  theta0 <- theta.post[i-1,]

  # mu's
  mu.mu = 0.5*(60*sigsq0 + sum(theta0))/(sigsq0 + 12)
  var.mu = sigsq0/(6 + 0.5*sigsq0)
  mu.post[i] = rnorm(1,mean = mu.mu, sd = sqrt(var.mu))

  # sigma-squareds
  beta.sigsq = 10 + 0.5*sum((t(y)-theta0)^2) + 0.5*sum((theta0 - mu0)^2)
  sigsq.post[i] = rinvgamma(1,shape = 122, scale = beta.sigsq)

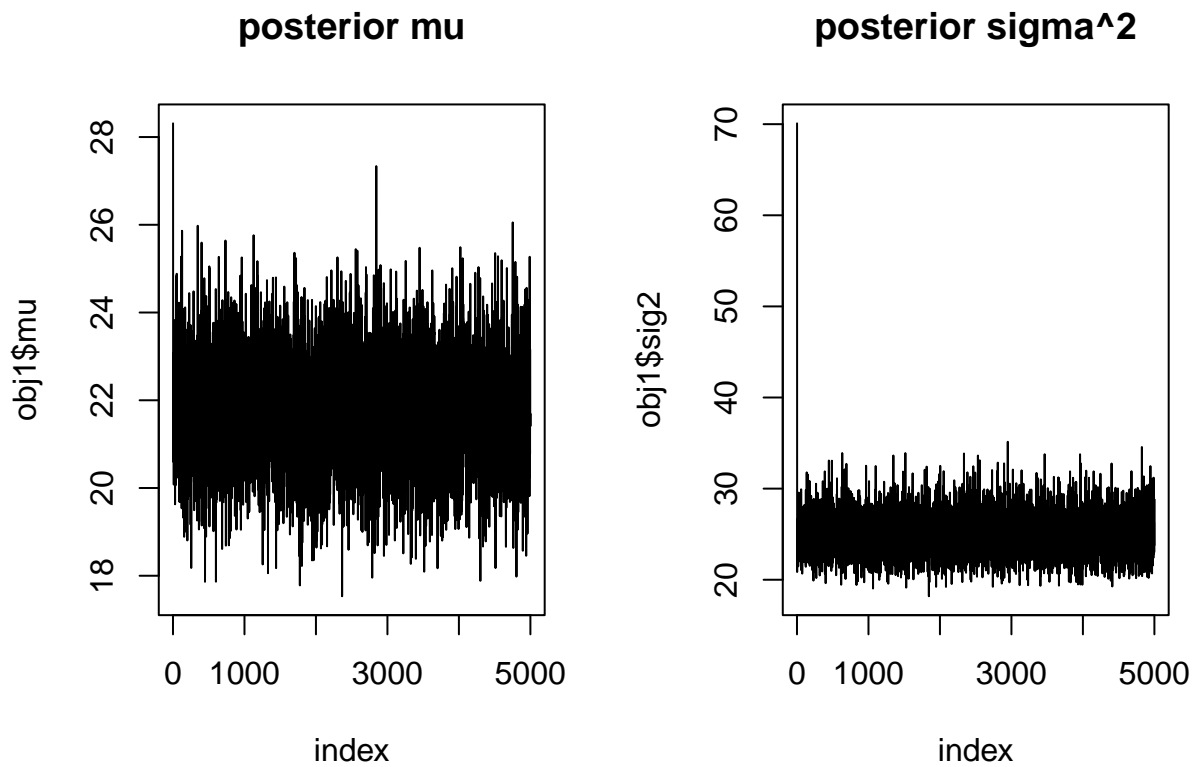
  # thetas
  for(j in 1:J){
    mu.theta = (mu0 + sum(y[,j]))/25
    sd.theta = sqrt(sigsq0/25)
    theta.post[i,j] = rnorm(1,mean = mu.theta, sd = sd.theta)
  }
}
return(list("theta"=theta.post,"mu"=mu.post,"sig2"=sigsq.post))
}

# put in some data
as.matrix(hear) -> y

GIBBS(y) -> obj1

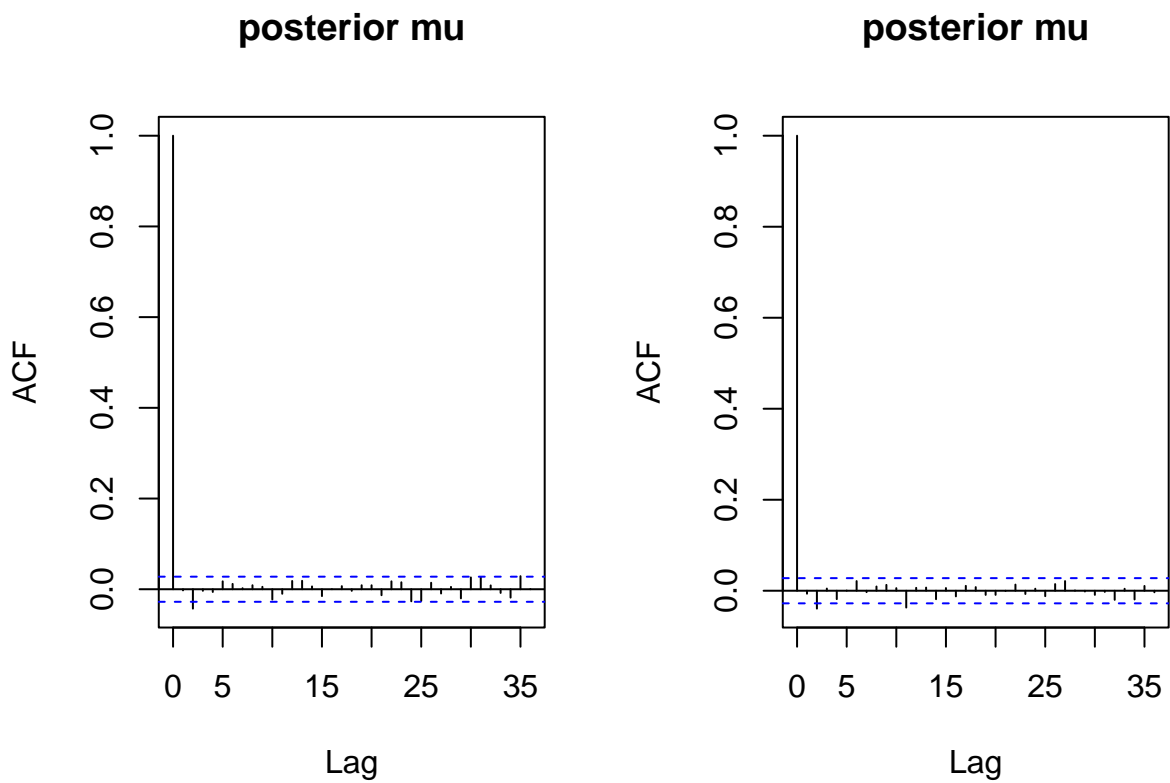
# visualizations
par(mfrow = c(1,2))
plot.ts(obj1$mu,main="posterior mu", xlab = "index")
plot.ts(obj1$sig2,main="posterior sigma^2", xlab = "index")

```



(e)

```
acf(obj1$mu,main="posterior mu")
acf(obj1$sig2,main="posterior mu")
```








no thinning needed for these parameters


```

#visualizing the four thetas
par(mfrow = c(1,2))
for(i in 1:4){
  plot.ts(obj1$theta[,i],xlab = "index")
  acf(obj1$theta[,i])
}

```



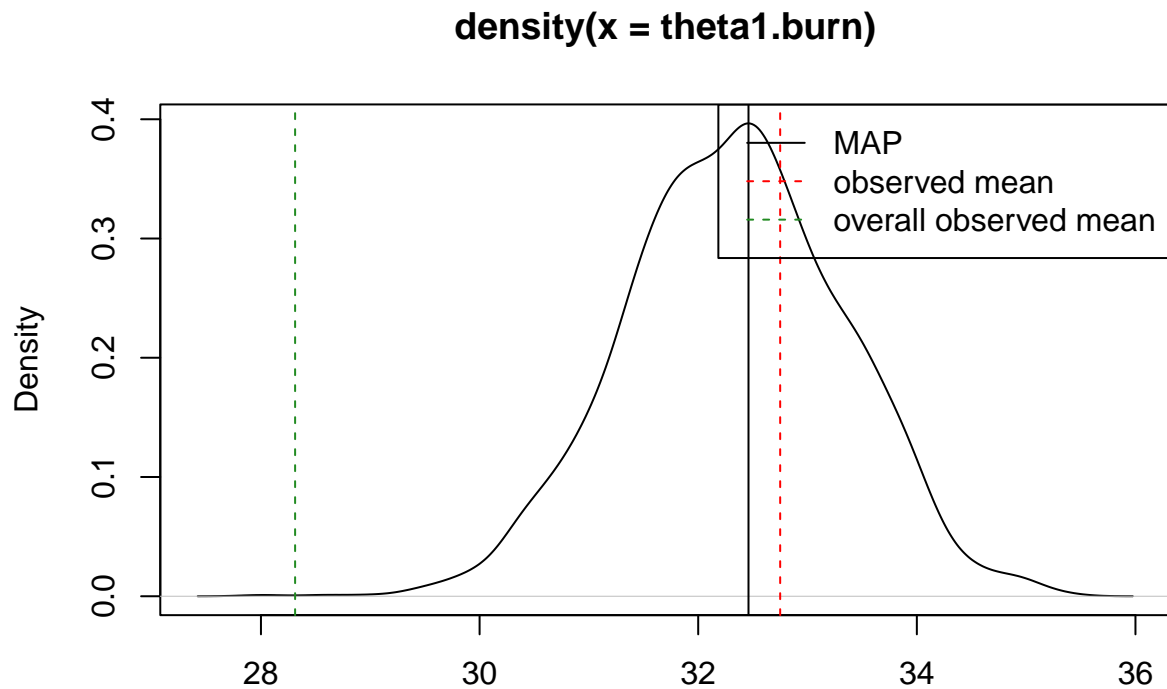




no thinning needed for *theta* either. Let's burn off the first 3000

```
# burning
k = 3001:5000
mu.burn = obj1$mu[k]
sig2.burn = obj1$sig2[k]
theta1.burn = obj1$theta[k,1]
theta2.burn = obj1$theta[k,2]
theta3.burn = obj1$theta[k,3]
theta4.burn = obj1$theta[k,4]
```

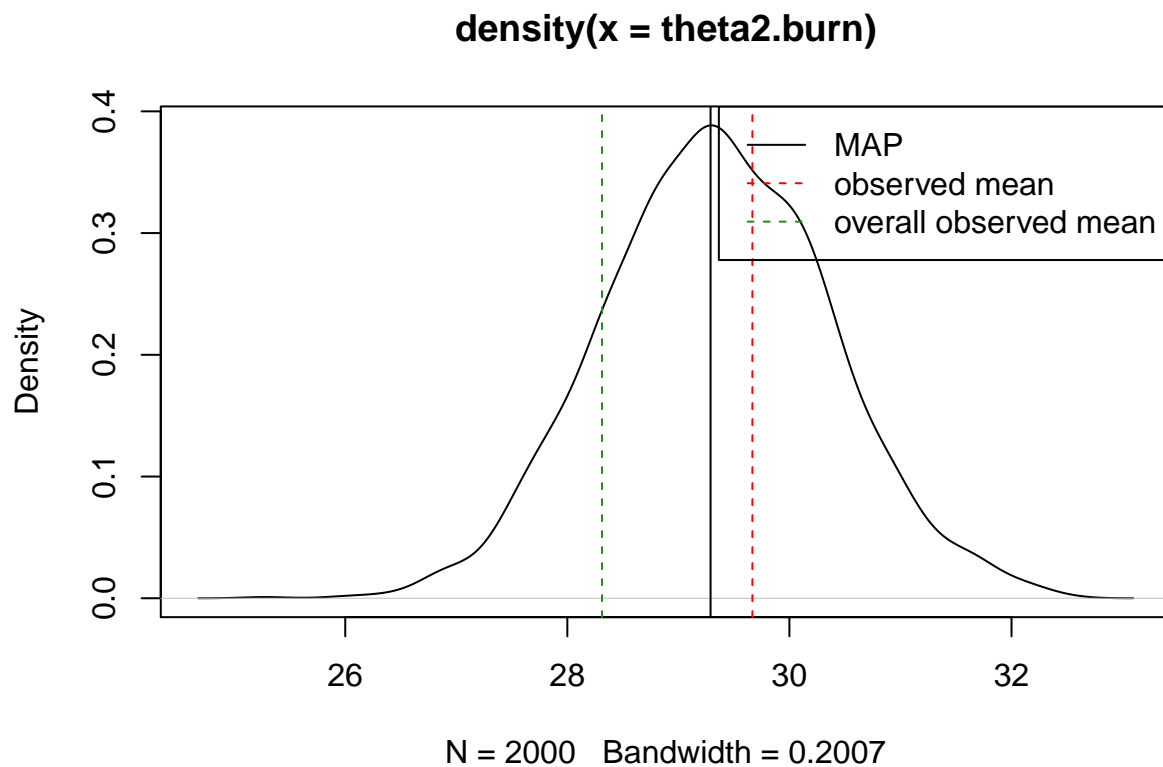
```
#par(mfrow=c(2,1))
MAP <- function(X){density(X)$x[density(X)$y == max(density(X)$y)]}
#theta1
dens = density(theta1.burn)
plot(dens)
abline(v = MAP(theta1.burn))
abline(v = mean(y[,1]), lty = 2, col = 'red')
abline(v = mean(y), lty = 2, col="forestgreen")
legend("topright",
      c("MAP", "observed mean", "overall observed mean"),
      lty = c(1,2,2),
      col = c("black", 'red', 'forestgreen'))
```

(f)


N = 2000 Bandwidth = 0.1973

```
#theta2
dens = density(theta2.burn)
plot(dens)
abline(v = MAP(theta2.burn))
abline(v = mean(y[,2]), lty = 2, col = 'red')
abline(v = mean(y), lty = 2, col = "forestgreen")
legend("topright",
      c("MAP", "observed mean", "overall observed mean"),
      lty = c(1, 2, 2),
      col = c("black", "red", "forestgreen"))
```



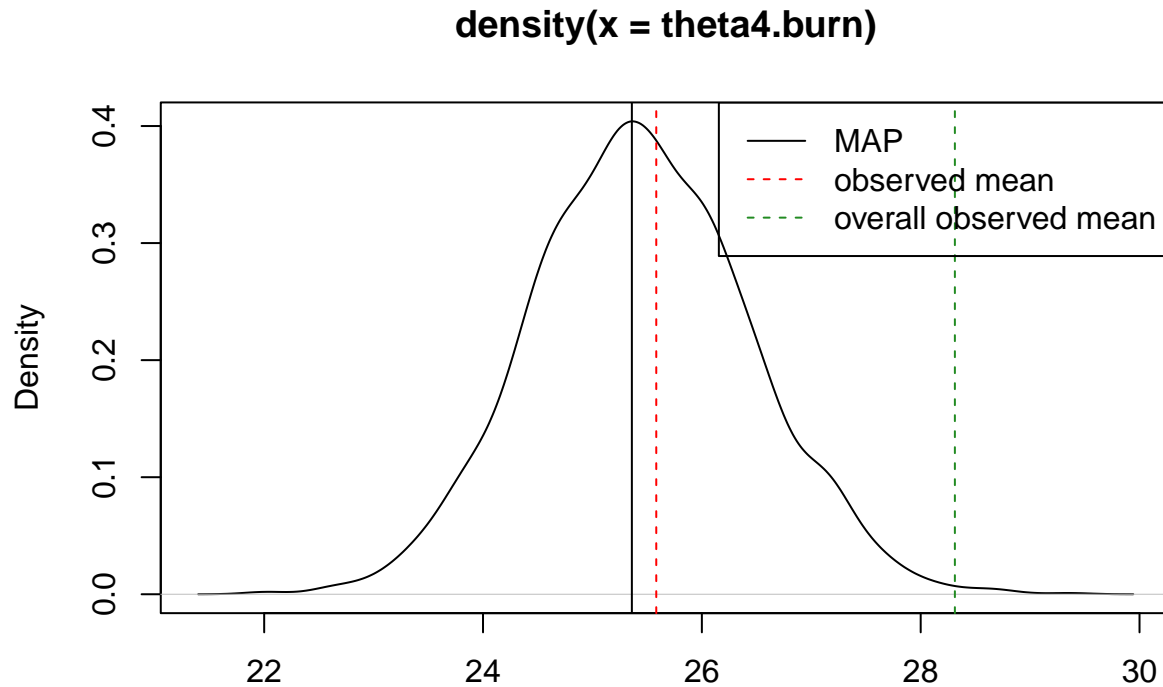
```
#theta3
dens = density(theta3.burn)
plot(dens)
abline(v = MAP(theta3.burn))
abline(v = mean(y[,3]), lty = 2, col = 'red')
abline(v = mean(y), lty = 2, col = "forestgreen")
legend("topright",
      c("MAP", "observed mean", "overall observed mean"),
      lty = c(1, 2, 2),
      col = c("black", "red", "forestgreen"))
```

density(x = theta3.burn)



N = 2000 Bandwidth = 0.2021

```
#theta4
dens = density(theta4.burn)
plot(dens)
abline(v = MAP(theta4.burn))
abline(v = mean(y[,4]), lty = 2, col = 'red')
abline(v = mean(y), lty = 2, col = "forestgreen")
legend("topright",
      c("MAP", "observed mean", "overall observed mean"),
      lty = c(1, 2, 2),
      col = c("black", 'red', 'forestgreen'))
```



N = 2000 Bandwidth = 0.1941

```
# credible intervals and point estimates
THET <- cbind(theta1.burn,theta2.burn,theta3.burn,theta4.burn)
row = matrix(0,4,3)
for(r in 1:4){
  row[r,] = c(quantile(THET[,r],0.025),MAP(THET[,r]),quantile(THET[,r],.975))
}
colnames(row) <- c("Lower","MAP","Upper")
rownames(row) <- c("theta_1","theta_2","theta_3","theta_4")
kable(row, caption = "95% credible interval and MAP")
```

Table 1: 95% credible interval and MAP

| | Lower | MAP | Upper |
|---------|----------|----------|----------|
| theta_1 | 30.35318 | 32.46017 | 34.17283 |
| theta_2 | 27.38634 | 29.28986 | 31.38136 |
| theta_3 | 23.11685 | 25.19115 | 27.07401 |
| theta_4 | 23.55233 | 25.36050 | 27.41182 |

It seems that the overall observed mean is always farther from the point estimate than the overall mean so there is clearly effect that being in a different list has; being in a different list can imply that there is there is a different level of difficulty that each list will experience.


```
# cred int
cred.map = rbind(c(quantile(sig2.burn,.025),MAP(sig2.burn),quantile(sig2.burn,.975)),
  c(quantile(mu.burn,.025),MAP(mu.burn),quantile(mu.burn,.975)))
rownames(cred.map) = c("sigma.squared", "mu")
colnames(cred.map) = c("Lower","MAP","Upper")
kable(cred.map, caption = "95% cred. interval and MAP")
```

(g)

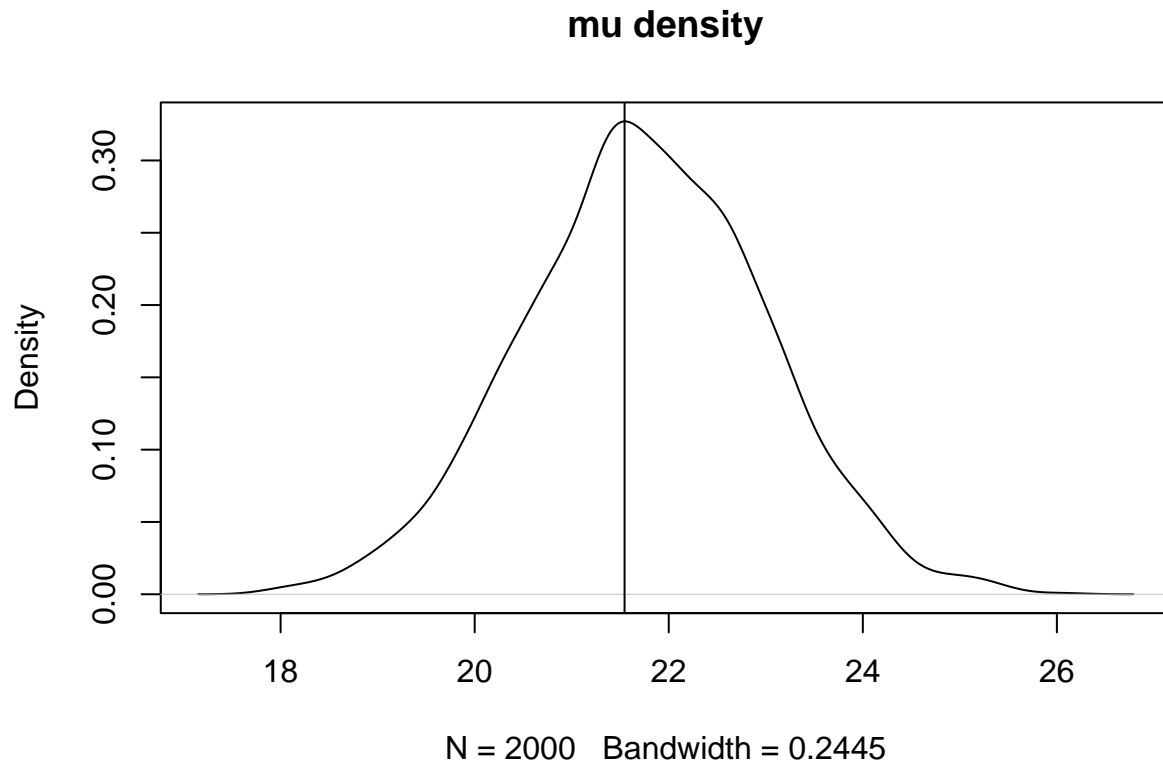
Table 2: 95% cred. interval and MAP

| | Lower | MAP | Upper |
|---------------|----------|----------|----------|
| sigma.squared | 20.77128 | 24.36342 | 30.11671 |
| mu | 19.30405 | 21.54458 | 24.17287 |

```
# density maps
#sigma-squared
dens = density(sig2.burn)
plot(dens, main = "sigma^2 density")
abline(v = MAP(sig2.burn))
```




```
#mu
dens = density(mu.burn)
plot(dens, main = "mu density")
abline(v = MAP(mu.burn))
```



Standard deviation for Y is around 5 for and the overall population mean of Y of the scores is around 22 which seems to make sense for list 3 and 4.

Problem 2

```
contr$p.hat <- contr$y/contr$N
plot(contr$p.hat,contr$W)
```




```
cor(contr$W, contr$p.hat)
```

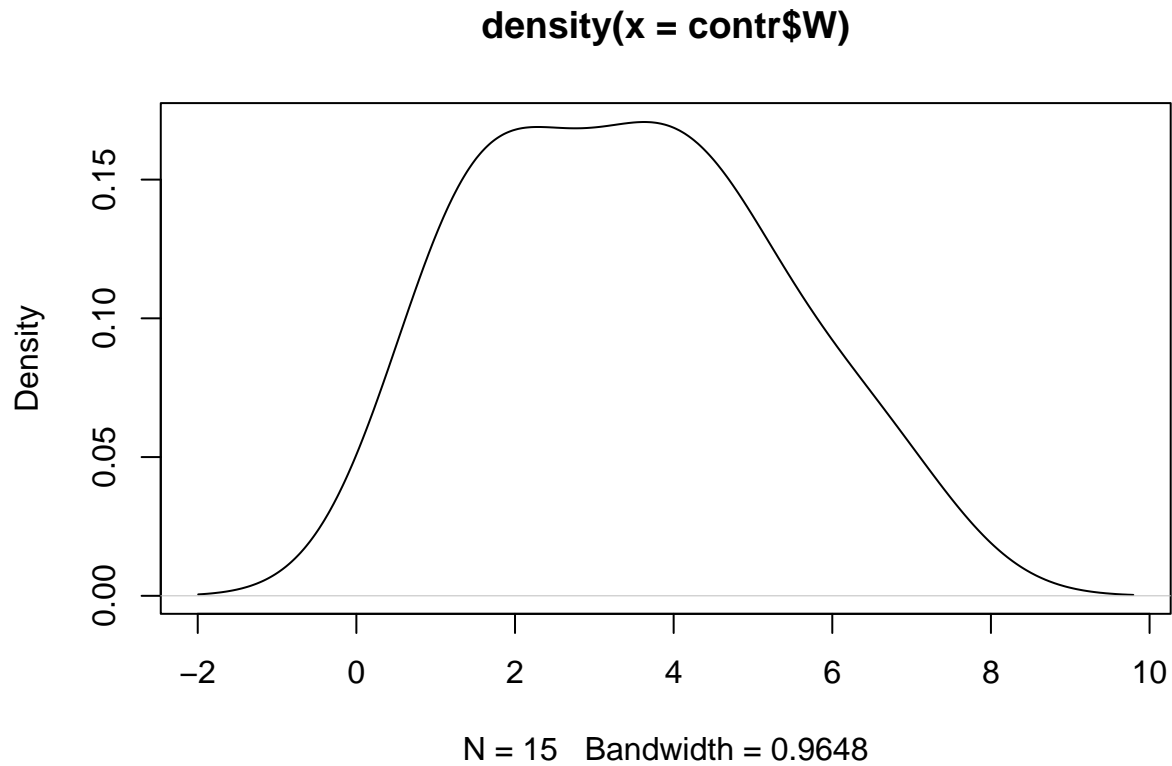
```
## [1] 0.4539023
```

There seems to be a slight positive correlation between proportion of women who use contraception and the education level but it's not too impressive (~ 0.4539). For simplicity sake, I have added a \hat{p} column.

```
plot(density(contr$p.hat))
```



```
plot(density(contr$W))
```



```
# define bayesians funcs
Prior <- function(alpha,beta){
  (alpha + beta)^(-5/2)
}

LLh <- function(theta,alpha,beta){
  sum(log(dbeta(theta, alpha, beta)))
}

Proposal <- function(alpha,beta){
  1/(alpha*beta)
}

rProposal <- function(n,mean,cov){
  rmvnorm(n,mean,cov)
}

# The Algorithm
Algo <- function(y,N,B,S){
  set.seed(538)
  # initials
  n = length(y)
  accept = 0
  alpha.post <- beta.post <- numeric()
  theta.post2 = matrix(0,ncol = B, nrow = n)

  # intial variables
  alpha0 = 4.1112
```



```

beta0 = 7
theta0 = numeric(length = n)

# LOOP
for(b in 1:B){
  # Gibbs step for thetas
  for(j in 1:n){
    shp1 = alpha0 + y[j]
    shp2 = beta0 + N[j] - y[j]
    theta0[j] = rbeta(1,shp1,shp2)
  }
  # MH step for alpha/beta
  phi.star = rProposal(1, c(log(alpha0),log(beta0)), 1*S)
  alpha.star = exp(phi.star[1])
  beta.star = exp(phi.star[2])

  r = exp(
    LLh(theta0, alpha.star, beta.star)
    + log(Prior(alpha.star, beta.star))
    + log(Proposal(alpha0, beta0))
    - LLh(theta0, alpha0, beta0)
    - log(Prior(alpha0, beta0))
    - log(Proposal(alpha.star, beta.star))
  )

  # logic selection
  if(runif(1) < min(1,r)){
    alpha0 = alpha.star
    beta0 = beta.star
    accept = accept + 1
  }
  # drop off
  alpha.post[b] = alpha0
  beta.post[b] = beta0
  theta.post2[,b] = theta0
}

# output
return(list("alpha"=alpha.post,"beta"=beta.post,"theta"=theta.post2,"AR"=accept/B))
}

```

```

# run the algorithm
Stune = diag(2)
contr$y -> yin
contr$N -> Nin
Algo(yin,Nin,50000,Stune) -> obj2
obj2$AR

```

(b)

```
## [1] 0.11308
```

```

# retune with posterior
a11 = var(obj2$alpha)

```

```

a12 <- cov(obj2$alpha,obj2$beta)
a22 <- var(obj2$beta)
Stune = matrix(c(a11,a12,a12,a22),2,2, byrow = T)
Stune

```

```

##           [,1]      [,2]
## [1,] 0.5585132 0.850026
## [2,] 0.8500260 1.821991

```

```

# rerun algorithm

```

```

Algo(yin,Nin,50000,Stune) -> obj2
obj2$AR

```

```

## [1] 0.18498

```

```

par(mfrow = c(1,2))

```

```

# visualizations

```

```

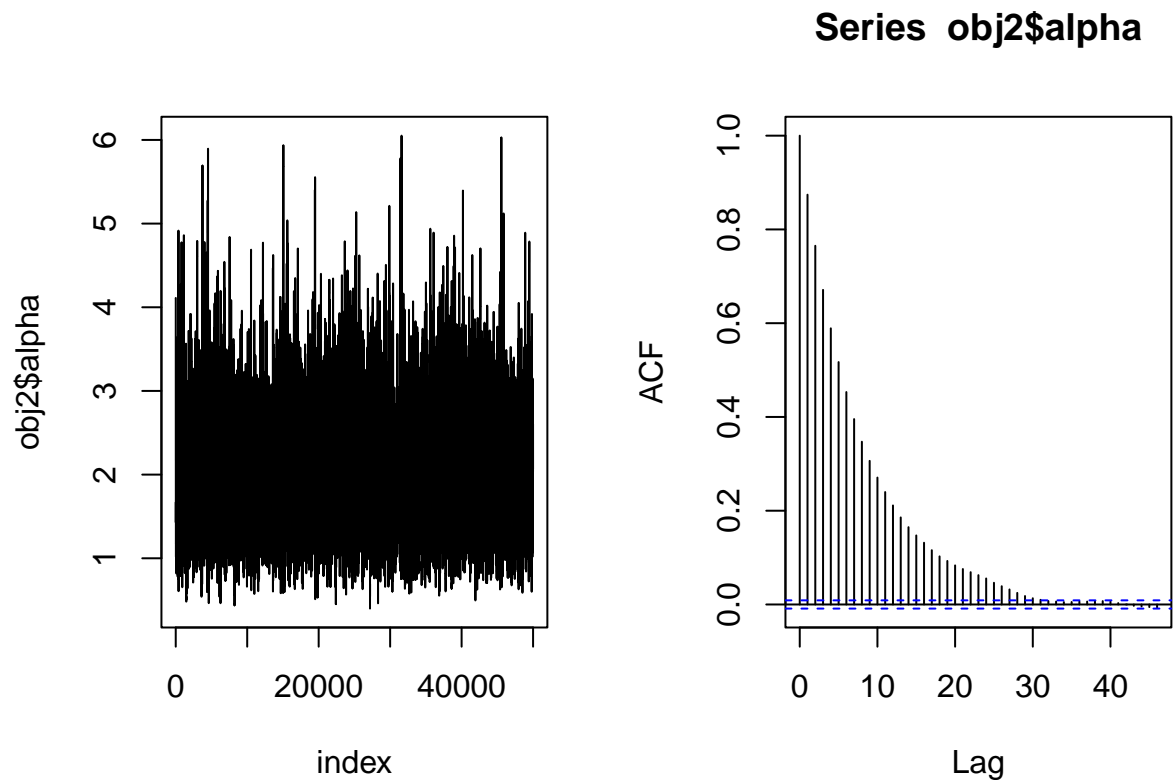
plot.ts(obj2$alpha,xlab = "index")

```

```

acf(obj2$alpha)



```



```

plot.ts(obj2$beta,xlab = "index")
acf(obj2$beta)



```



```
B = length(obj2$alpha)
```



```
burn = 20000
1:B*40 -> k
k = k[burn < k & k < B]
alpha.thin = obj2$alpha[k]
beta.thin = obj2$beta[k]
# burned/thinned visualizations
par(mfrow = c(1,2))
plot.ts(alpha.thin,xlab = "index")
acf(alpha.thin)
```

Series alpha.thin




```
plot.ts(beta.thin,xlab = "index")  
acf(beta.thin)
```

Series beta.thin



```
# visuals  
dens = density(alpha.thin)
```


```
plot(dens)
abline(v = MAP(alpha.thin))
```



N = 749 Bandwidth = 0.1727

```
dens = density(beta.thin)
plot(dens)
abline(v = MAP(beta.thin))
```


density(x = beta.thin)



N = 749 Bandwidth = 0.3097

```
ev = alpha.thin/(alpha.thin+beta.thin)
dens = density(ev)
plot(dens, main = "E[theta]")
abline(v = MAP(ev))
```

E[theta]



N = 749 Bandwidth = 0.01222

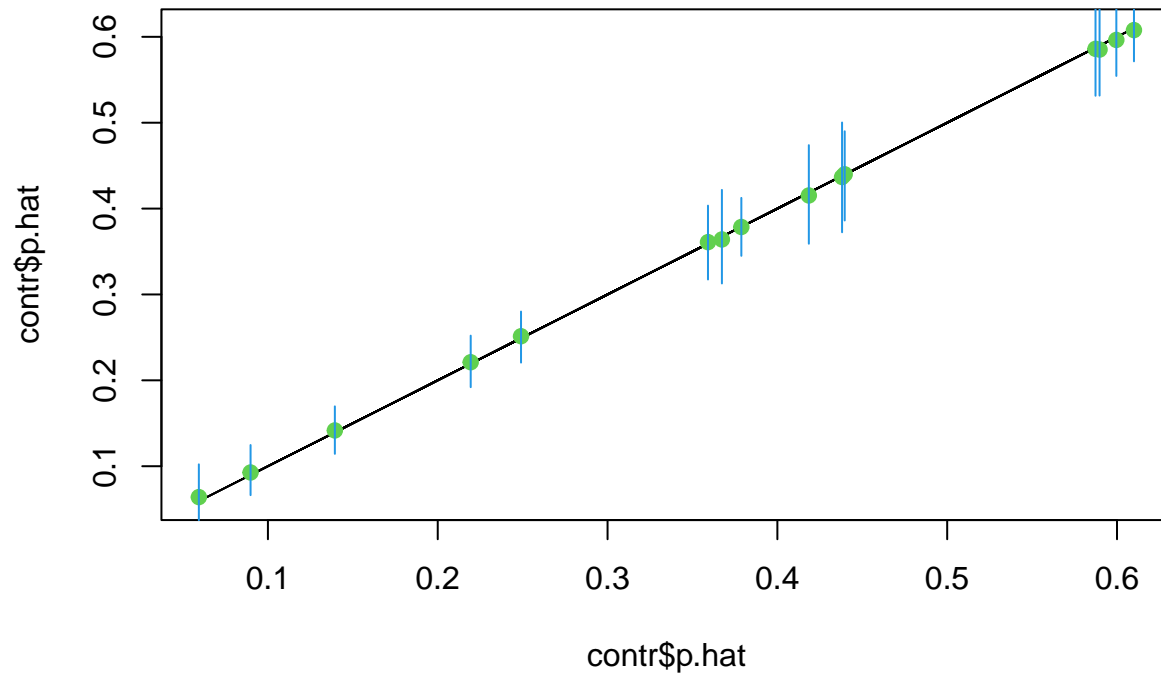
```
# table
cred.map = rbind(
  c(quantile(alpha.thin,0.025),MAP(alpha.thin),quantile(alpha.thin,.975)),
  c(quantile(alpha.thin,0.025),MAP(beta.thin),quantile(beta.thin,.975)),
  c(quantile(ev,0.025),MAP(ev),quantile(ev,.975))
)
rownames(cred.map) = c("alpha", "beta", "E[theta]")
colnames(cred.map) = c("Lower","MAP","Upper")
kable(cred.map, caption = "95% cred. interval and MAP")
```

Table 3: 95% cred. interval and MAP

| | Lower | MAP | Upper |
|----------|-----------|-----------|-----------|
| alpha | 0.8465403 | 1.7861238 | 3.4966208 |
| beta | 0.8465403 | 2.7132091 | 6.1132207 |
| E[theta] | 0.2796714 | 0.3710374 | 0.4729117 |

I think we may be able to conclude that there is justification in using this model as our credible interval for $E[\theta]$ contains the $\bar{\theta}$ value of 0.4539023.

```
theta.thin = obj2$theta[,k]
#----- Displaying the theta's
CI = matrix(0, ncol = 3, nrow = 15)
for (j in 1:15) {
  CI[j, ] = quantile(theta.thin[j, ], probs = c(0.025,0.5, 0.975))}
plot(contr$p.hat, contr$p.hat, type = "l")
points(contr$p.hat, CI[, 2], pch = 19, col = 3)
for (j in 1:15) {
  points(c(contr$p.hat[j], contr$p.hat[j]), c(CI[j, 1], CI[j,3]), type = "l", col = 4)}
```



```

cred.map = matrix(0,15,3)
for(i in 1:15){
  v = as.vector(quantile(theta.thin[i,],probs = c(0.025,.5,.975)))
  cred.map[i,] =v
}
colnames(cred.map) = c("Lower", "MAP", "Upper")
kable(cred.map)

```

| Lower | MAP | Upper |
|-----------|-----------|-----------|
| 0.2206381 | 0.2513666 | 0.2800969 |
| 0.1920471 | 0.2210838 | 0.2520631 |
| 0.3449612 | 0.3785759 | 0.4124509 |
| 0.5714656 | 0.6078265 | 0.6445202 |
| 0.1143959 | 0.1416608 | 0.1697049 |
| 0.5544570 | 0.5963938 | 0.6383091 |
| 0.3174360 | 0.3609464 | 0.4034365 |
| 0.0662946 | 0.0926787 | 0.1246903 |
| 0.3861322 | 0.4399413 | 0.4900302 |
| 0.5313950 | 0.5860343 | 0.6353225 |
| 0.3590568 | 0.4152829 | 0.4738853 |
| 0.5315867 | 0.5851860 | 0.6386907 |
| 0.3128118 | 0.3640594 | 0.4217566 |
| 0.0370283 | 0.0640338 | 0.1022062 |
| 0.3724190 | 0.4367750 | 0.5002366 |

```

kable(contr$W)

```

```

—
x
—
1.2
3.8
1.6
4.5
2.7
3.7
2.1
0.9
1.4
4.7
6.1
5.3
2.7
3.9
6.9
—

```

Looking at the graph above, we can conclude that our posterior θ 's imply that we need to account for cluster differences. Thus the hierarchical model is appropriate.

analyzing the literature

I believe that the main objective of the authors was to show that the model on medicare.gov was it rather misleading in predicting their mortality rate (of heart attack victims); it was important for them to show there's a more efficient way of estimating a mortality rate given the various attributes of the medical centers at hand. The issues that the statisticians ran into considering the hierarchical model being used for "hospital compare": the model was inaccurately following the data when considering lower volume medical centers. Apparently the previous model was making assumptions about spread and expected value that, not accounting for attributes such as size, were not fitting to the data. Smaller volume hospitals will always have less data than bigger facilities; so spread, or variance, is going to be naturally different when throwing them into a pool where data sets from larger hospital swim. The mortality rate for these lower volume hospitals was "shrinking" down to the national average. Imagine this situation person at risk for heart attack comparing hospitals whose mortality rates are being underestimated. Overall, the authors found that considering more hospital attributes (as well as patient risk factors and their relation to the hospital) is what makes the better model. Personally, I think this could be detrimental to public policy as policy makers who are unaware of the nuances of the topic at hand can be misled to make poor decisions because of a faulty statistical model; we must be sure our assumptions align with our observations. Three small improve their mortality rate they received more funding but miss out on it because of politicians who are lead by the numbers of "hospital compare."